

Notes

Machine Learning by Andrew Ng on Coursera

Sparsh Jain

November 11, 2020

Contents

1	Introduction	5
1.1	Supervised Learning	5
1.1.1	Regression Problem	5
1.1.2	Classification Problems	5
1.2	Unsupervised Learning	5
2	Linear Regression with One Variable	6
2.1	Notations	6
2.2	Supervised Learning	6
2.3	Gradient Descent	7
2.4	Gradient Descent for Linear Regression	7
3	Linear Algebra	8
3.1	Matrix	8
3.2	Vector	8
3.3	Addition and Scalar Multiplication	9
3.4	Matrix Matrix Multiplication	9
3.5	Inverse and Transpose	9
4	Linear Regression with Multiple Variables	11
4.1	Notations	11
4.2	Hypothesis	11
4.3	Gradient Descent	12
4.3.1	Feature Scaling	12
4.3.2	Mean Normalization	13
4.3.3	Learning Rate	13
4.4	Features and Polynomial Regression	13
4.4.1	Features	13
4.4.2	Polynomial Regression	13
4.5	Normal Equation	14
4.5.1	Non Invertibility of $X^T X$	14

5	Octave Tutorial	15
6	Classification	16
6.1	Logistic Regression	16
6.2	Decision Boundary	17
6.3	Cost Function	17
6.4	Advanced Optimization:	19
6.5	Multi-Class Classification	20
6.5.1	One vs All	20
7	Regularization	21
7.1	Problem of Overfitting	21
7.2	Addressing Overfitting	21
7.3	Cost Function	22
7.4	Regularized Linear Regression	22
7.5	Regularized Logistic Regression	24
8	Neural Networks	25
8.1	Non-Linear Hypothesis	25
8.2	Neural Networks	25
8.3	Model Representation	25
8.4	Notations	26
8.5	Forward Propagation	27
8.6	Multi-Class Classification	28
9	Back Propagation	29
9.1	Notations	29
9.2	Cost Function	30
9.3	Backpropagation Algorithm	30
9.3.1	Gradient Computation	30
9.4	Implementation Note	32
9.4.1	Unrolling Parameters	32
9.4.2	Learning Algorithm	33
9.5	Gradient Checking	33
9.5.1	Numerical Estimation of gradients	33
9.5.2	Implementation Note:	35
9.6	Random Initialization	35
9.7	Putting it all together	36

10 Applying Machine Learning	38
10.1 Deciding what to try next	38
10.1.1 Debugging a Learning Algorithm	38
10.1.2 Machine Learning Diagnostics	38
10.2 Evaluating your Hypothesis	39
10.2.1 The Standard Way	39
10.2.2 Train/Test Procedure	39
10.3 Model Selection	39
10.3.1 Evaluating your hypothesis	40
10.4 Bias vs Variance	41
10.4.1 Effect of regularization	41
10.5 Learning Curves	42
10.5.1 High Bias	42
10.5.2 High Variance	43
10.6 Deciding what to try next (Revisited)	43
10.6.1 Neural Networks	43
11 Machine Learning System Design	44
11.1 Prioritizing what to work on: Spam Classification Example . .	44
11.1.1 Time Management?	45
11.2 Error Analysis	45
11.2.1 Recommended approach	45
11.2.2 Numerical Evaluation	45
11.3 Skewed Classes	46
11.3.1 Precision/Recall	46
11.3.2 Trade-Off	47
11.3.3 Working with Large Data	47
12 Support Vector Machines	48
12.1 Optimization Objective	48
12.1.1 Alternative view of logistic regression	48
12.1.2 SVM Hypothesis	49
12.2 Large Margin	50
12.2.1 Intuition	50
12.2.2 Math	50
12.3 Kernels	51
12.3.1 Non-linear Decision Boundary	51
12.3.2 Choosing the Landmarks	52
12.3.3 SVM Parameters:	53
12.4 Using SVM	53
12.4.1 Multi-Class Classification	54

12.4.2 Logistic Regression vs SVM	54
---	----

Chapter 1

Introduction

Machine learning (task, experience, performance) can be classified into *Supervised* and *Unsupervised* learning.

1.1 Supervised Learning

Supervised learning can be basically classified into *Regression* and *Classification* problems.

1.1.1 Regression Problem

Regression problems work loosely on continuous range of outputs.

1.1.2 Classification Problems

Classification problems work loosely on discrete range of outputs.

1.2 Unsupervised Learning

An example is *Clustering Problem*.

Check Lecture1.pdf for more details.

Chapter 2

Linear Regression with One Variable

2.1 Notations

m = number of training examples

x 's = 'input' variables / features

y 's = 'output' variables / 'target' variables

(x, y) = single training example

$(x^{(i)}, y^{(i)}) = i^{th}$ example

2.2 Supervised Learning

We have a data set (*Training Set*).

Training Set \rightarrow Learning Algorithm $\rightarrow h$ (*hypothesis*, a function $X \rightarrow Y$)

To Represent h

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Cost

$$\underset{\theta_0, \theta_1}{\text{minimize}} \frac{1}{2m} \sum_1^m (h_{\theta}(x) - y)^2$$

Cost Function

Squared Error Cost Function

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_1^m (h_{\theta}(x) - y)^2$$

$$\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$$

2.3 Gradient Descent

Finds local optimum:

1. Start with some value
2. Get closer to optimum

Algorithm

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \quad \forall j$$

where α = learning rate

Important!

Simultaneous Update!

$$\begin{aligned} temp_j &:= \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \quad \forall j \\ \theta_j &:= temp_j \quad \forall j \end{aligned}$$

2.4 Gradient Descent for Linear Regression

Cost function for linear regression is convex!

Batch Gradient Descent: Each step of gradient descent uses all training examples.

Check Lecture2.pdf for more details.

Chapter 3

Linear Algebra

3.1 Matrix

Rectangular array of numbers:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

Dimension of the matrix: #rows x #cols (2 x 3)

Elements of the matrix:

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

A_{ij} = “ i, j entry” in the i^{th} row, j^{th} col

3.2 Vector

An $n \times 1$ matrix.

$$y = \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix}$$

$$y_i = i^{th} \text{ element}$$

Note: Uppercase for matrices, lowercase for vectors.

3.3 Addition and Scalar Multiplication

Add/Subtract (element by element) matrices of same dimension only!

Multiply/Divide (all elements) a matrix by scalar!

3.4 Matrix Matrix Multiplication

$m \times n$ matrix multiplied by $n \times o$ matrix gives a $m \times o$ matrix.

Properties

1. Matrix Multiplication is *not* Commutative.
2. Matrix Multiplication is Associative.
3. *Identity Matrix (I)*: 1's along diagonal, 0's everywhere else in an $n \times n$ matrix. $AI = IA = A$.

3.5 Inverse and Transpose

Inverse

Only square ($n \times n$) matrices *may* have an inverse.

$$AA^{-1} = A^{-1}A = I$$

Matrices that don't have an inverse are *singular* or *degenerate* matrices.

Transpose

Let A be an $m \times n$ matrix and let $B = A^T$, then

$$B_{ij} = A_{ji}$$

Example:

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$
$$B = A^T = \begin{bmatrix} 1 & 4 \\ 2 & 5 \\ 3 & 6 \end{bmatrix}$$

Check Lecture3.pdf for more details.

Chapter 4

Linear Regression with Multiple Variables

4.1 Notations

n = number of features

$x^{(i)}$ = input (features) of i^{th} training example

$x_j^{(i)}$ = value of feature j of i^{th} training example

4.2 Hypothesis

Previously:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Now:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

For convinience, define $x_0 = 1$. So

$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$h_{\theta}(x) = \theta^T x$$

4.3 Gradient Descent

Hypothesis : $h_{\theta}(x) = \theta^T x$

Parameters : θ

$$= \theta_0 x_0 + \theta_1 x_1 + \cdots + \theta_n x_n$$

$$= \theta_0, \theta_1, \dots, \theta_n$$

Cost Function : $J(\theta) = J(\theta_0, \theta_1, \dots, \theta_n)$

$$= \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Gradient Descent :

Repeat{

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

$$= \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1, \dots, \theta_n)$$

$$= \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

} (simultaneously update $\forall j = 0, 1, \dots, n$)

4.3.1 Feature Scaling

Idea: Make sure features are on a similar scale.

Get every feature into approximately a $-1 \leq x_i \leq 1$ range.

4.3.2 Mean Normalization

Replace x_i with $x_i - \mu_i$ to make features have approximately zero mean (Do not apply to $x_0 = 1$).

General Rule

$$x_i \leftarrow \frac{x_i - \mu_i}{S_i}$$

where

μ_i = average value of x_i

S_i = range (max - min) *or*

= σ (standard deviation)

4.3.3 Learning Rate

$J(\theta)$ should decrease after every iteration. #iterations vary a lot.

Example *Automatic Convergence Test*: Declare convergence if $J(\theta)$ decreases by less than ϵ (say 10^{-3}) in one iteration.

If $J(\theta)$ increases, use smaller α . Too small α means slow convergence.

To choose α , try ..., 0.001, 0.003, 0.01, 0.03, 0.1, 0.3, 1, ...

4.4 Features and Polynomial Regression

4.4.1 Features

Get an insight in your problem and choose better features (may even combine/separate features).

Ex: size = length \rightarrow breadth.

4.4.2 Polynomial Regression

Ex:

$$x_1 = size$$

$$x_2 = size^2$$

$$x_3 = size^3$$

4.5 Normal Equation

Solve for θ analytically!

$$\begin{aligned}x^{(i)} &= \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix} && \in \mathbb{R}^{n+1} \\X &= \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ \vdots \\ (x^{(m)})^T \end{bmatrix} && \in \mathbb{R}^{m \times (n+1)} \\&= \begin{bmatrix} x_0 & x_1 & \dots & x_n \end{bmatrix} && \in \mathbb{R}^{m \times (n+1)} \\y &= \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix} && \in \mathbb{R}^m \\\theta &= (X^T X)^{-1} X^T y\end{aligned}$$

Inverse of a matrix grows as $O(n^3)$, use wisely.

4.5.1 Non Invertibility of $X^T X$

Use 'pinv' function in Octave (pseudo-inverse) instead of 'inv' function (inverse).

If $X^T X$ is non-invertible, common causes are

1. Redundant features (linearly dependent)
2. Too many features ($m \leq n$). In this case, delete some features or use *regularization*

Check Lecture4.pdf for more details.

Chapter 5

Octave Tutorial

Check Lecture5.pdf for more details.

Chapter 6

Classification

Classify into categories (binary or multiple).

6.1 Logistic Regression

$$\begin{aligned} 0 &\leq h_{\theta}(x) \leq 1 \\ h_{\theta}(x) &= g(\theta^T x) \\ g(z) &= \frac{1}{1 + e^{-z}} \quad g \text{ is called a } \textit{sigmoid} \text{ function or a } \textit{logistic} \text{ function.} \\ h_{\theta}(x) &= \frac{1}{1 + e^{\theta^T x}} \end{aligned}$$

Interpretation of Hypothesis Output

$h_{\theta}(x)$ = estimated probability that $y = 1$ on input x

$h_{\theta}(x) = P(y = 1|x; \theta)$ = probability that $y = 1$, given x , parameterized by θ

$$P(y = 0|x; \theta) + P(y = 1|x; \theta) = 1$$

6.2 Decision Boundary

$$\begin{array}{ll} \text{Predict: } y = 1 & \text{if } h_{\theta}(x) \geq 0.5 \\ & (\theta^T x \geq 0) \\ \text{Predict: } y = 0 & \text{if } h_{\theta}(x) < 0.5 \\ & (\theta^T x < 0) \\ \theta^T x = 0 & \text{is the } \textit{decision boundary}. \end{array}$$

Non-linear Decision Boundaries

Use same technique as polynomial regression for features.

6.3 Cost Function

Training Set : $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$x_0 = 1$$

$$y \in \{0, 1\}$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

How to choose parameter θ ?

Linear Regression:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

$$\text{Cost}(h_{\theta}(x), y) = \frac{1}{2} (h_{\theta}(x) - y)^2$$

Logistic Regression:

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & y = 1 \\ -\log(1 - h_{\theta}(x)) & y = 0 \end{cases}$$

$$\text{Cost}(h_{\theta}(x), y) = 0 \text{ if } h_{\theta}(x) = y$$

$$\text{Cost}(h_{\theta}(x), y) \rightarrow \inf \text{ if } y = 0 \text{ and } h_{\theta}(x) \rightarrow 1$$

$$\text{Cost}(h_{\theta}(x), y) \rightarrow \inf \text{ if } y = 1 \text{ and } h_{\theta}(x) \rightarrow 0$$

Note: $y = 0$ or 1 always.

$$\text{Cost}(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$$

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[\sum_{i=1}^m (y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))) \right] \end{aligned}$$

To fit parameters θ :

$$\underset{\theta}{\text{minimize}} J(\theta)$$

To make a prediction given a new x :

$$\text{Output } h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Gradient Descent:

Simultaneously update all θ_j

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Plug in the derivative

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Don't forget feature scaling!

6.4 Advanced Optimization:

Something better than gradient descent:

1. Conjugate Gradient
2. BFGS
3. L-BFGS

Advantages:

1. No need to manually pick α
2. Often faster than gradient descent

Disadvantages:

1. More complex

Use libraries! Beware of bad implementations!

How to use: We first need to provide a function that evaluates the following two functions for a given input value of θ .

1. $J(\theta)$
2. $\frac{\partial}{\partial \theta_j} J(\theta)$

```

% We can write a single function that can return both of these:
function [jVal, gradient] = costFunction(theta)
    jVal = [...code to compute J(theta)...];
    gradient = [...code to compute derivative of J(theta)...];
end

```

Then we can use octave's *fminunc()* optimization algorithm along with the *optimset()* function that creates an object containing the options we want to send to *fminunc()*.

```

options = optimset('GradObj', 'on', 'MaxIter', 100);
initialTheta = zeros(2,1);
[optTheta, functionVal, exitFlag] = fminunc(@costFunction,
    initialTheta, options);

```

6.5 Multi-Class Classification

6.5.1 One vs All

Build a separate binary classifier $h_{\theta}^{(i)}(x)$ for each class against all other classes.

$$h_{\theta}^{(i)} = P(y = i|x; \theta) \quad \forall i$$

On a new input x , to make a prediction, pick the class i that maximizes $h_{\theta}^{(i)}(x)$

Check Lecture6.pdf for more details.

Chapter 7

Regularization

7.1 Problem of Overfitting

1. Underfitting (High Bias)
2. Right Fit
3. Overfitting (High Variance)

Overfitting: If we have too many features, the learned hypothesis may fit the training set very well $\left(J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h\theta(x^{(i)}) - y^{(i)})^2 \approx 0 \right)$, but fail to generalize to new examples (predict prices on new examples).

7.2 Addressing Overfitting

Options:

1. Reduce the number of features
 - Manually select which features to keep
 - Model selection algorithm (later)
2. Regularization
 - Keep all the features, but reduce the magnitude/values of parameters θ_j
 - Works well when we have a lot of features, each of which contributes a bit to predicting y

7.3 Cost Function

Say our overfitting hypothesis is $h_\theta(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$. Suppose, we penalize and make θ_3, θ_4 very small

$$\underset{\theta}{\text{minimize}} \left(\frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2 + 1000\theta_3^2 + 1000\theta_4^2 \right)$$

Regularization

Small values for parameters $\theta_0, \theta_1, \dots, \theta_n$.

- *Simpler* hypothesis
- Less prone to overfitting

Which parameters to penalize?

- Features: x_1, x_2, \dots, x_{100}
- Parameters: $\theta_0, \theta_1, \theta_2, \dots, \theta_{100}$

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

Note: The convention, we don't regularize θ_0 but it doesn't make very much difference.

λ here is *regularization parameter*.

What if λ is set too high (say 10^{10})? Underfitting!

7.4 Regularized Linear Regression

Updated $J(\theta)$:

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$
$$\underset{\theta}{\text{minimize}} J(\theta)$$

Gradient Descent:

Repeat{

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m \left((h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{m} \theta_j \right] \quad \forall j \in \{1, 2, \dots, n\}$$

$$\equiv \theta_j := \theta_j \left(1 - \alpha \frac{\lambda}{m}\right) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}

Normal Equation:

$$X = \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ \vdots \\ (x^{(m)})^T \end{bmatrix} \in \mathbb{R}^{m \times (n+1)}$$

$$y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix} \in \mathbb{R}^m$$

$$\theta = \left(X^T X + \lambda \begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ 0 & 0 & 0 & \ddots & 0 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix} \right)^{-1} \begin{pmatrix} \in \mathbb{R}^{(n+1) \times (n+1)} \end{pmatrix} - X^T y \in \mathbb{R}^{n+1}$$

Non-InvertibilitySuppose $m \leq n$,

$$\theta = (X^T X)^{-1} X^T y$$

If $\lambda > 0$,

$$\theta = \left(X^T X + \lambda \begin{bmatrix} 0 & & & \\ & 1 & & \\ & & 1 & \\ & & & \ddots \\ & & & & 1 \end{bmatrix} \right)^{-1} - X^T y$$

Not a problem!

7.5 Regularized Logistic Regression

Cost Function:

$$J(\theta) = - \left[\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Gradient Descent:

Repeat{

$$\begin{aligned} \theta_0 &:= \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)} \\ \theta_j &:= \theta_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}) + \frac{\lambda}{m} \theta_j \right] \quad \forall j \in \{1, 2, \dots, n\} \\ &\equiv \theta_j := \theta_j (1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \\ &\} \end{aligned}$$

Advanced Optimization:

```
function [jVal, gradient] = costFunction(theta)
    jVal = [...code to compute J(theta)...];
    gradient = [...code to compute derivative of J(theta)...];
```

Check Lecture7.pdf for more details.

Chapter 8

Neural Networks

8.1 Non-Linear Hypothesis

If #features is high, hypothesis could have extremely high #terms. Hence the need for non-linear hypothesis.

8.2 Neural Networks

- *Origins:* Algorithms that try to mimic brain.
- Widely used in 80s and early 90s; diminished in late 90s.
- Resurgence: State-of-the-art technique for many applications.

8.3 Model Representation

Neuron Model: Logistic Unit

1. input layer
2. hidden layer / computation layer
3. output layer
4. activation function ($g()$)

parameters \equiv weights

8.4 Notations

$a_i^{(j)}$ = *activation* of unit i in layer j
 $\Theta^{(j)}$ = matrix of weights controlling
function mapping from layer j to
layer $j + 1$

Example:

layer 1 : $\{x_1, x_2, x_3\}$

layer 2 : $\{a_1^{(2)}, a_2^{(2)}, a_3^{(2)}\}$

layer 3 : $\{a_1^{(3)}\}$

$$\begin{aligned}a_1^{(2)} &= g \left(\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3 \right) \\a_2^{(2)} &= g \left(\Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3 \right) \\a_3^{(2)} &= g \left(\Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3 \right) \\h_{\Theta}(x) = a_1^{(3)} &= g \left(\Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)} \right)\end{aligned}$$

Note: If network has s_j units in layer j , and s_{j+1} units in layer $j + 1$, then

$$\Theta^{(j)} \in \mathbb{R}^{s_{j+1} \times (s_j + 1)}$$

8.5 Forward Propagation

$$z_1^{(2)} = \Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3$$

$$z_2^{(2)} = \Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3$$

$$z_3^{(2)} = \Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3$$

$$a_1^{(2)} = g(z_1^{(2)})$$

$$a_2^{(2)} = g(z_2^{(2)})$$

$$a_3^{(2)} = g(z_3^{(2)})$$

Vectorized Implementation

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix}, \quad x_0 = 1$$

$$z^{(2)} = \begin{bmatrix} z_1^{(2)} \\ z_2^{(2)} \\ z_3^{(2)} \end{bmatrix}$$

$$a^{(1)} = x \quad \in \mathbb{R}^4$$

$$z^{(2)} = \Theta^{(1)} a^{(1)} \quad \in \mathbb{R}^3$$

$$a^{(2)} = g(z^{(2)}) \quad \in \mathbb{R}^3$$

$$\text{Add } a_0^{(2)} = 1 \quad \rightarrow a^{(2)} \in \mathbb{R}^4$$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

$$h_{\Theta}(x) = a^{(3)} = g(z^{(3)})$$

Other network architectures are possible!

Non-linear classification example: XOR/XNOR

x_1, x_2 are binary (0 or 1)

$$\begin{aligned} y &= x_1 \text{ XOR } x_2 && \text{or} \\ &= x_1 \text{ XNOR } x_2 && \equiv \text{NOT}(x_1 \text{ XOR } x_2) \end{aligned}$$

¹ Simple example: AND

$$\begin{aligned} x_0 &= 0 \\ x_1, x_2 &\in \{0, 1\} \\ y &= x_1 \text{ AND } x_2 \\ \Theta_{10}^{(1)} &= -30 \\ \Theta_{11}^{(1)} &= 20 \\ \Theta_{12}^{(1)} &= 20 \\ h_{\Theta}(x) &= g(-30 + 20x_1 + 20x_2) \end{aligned}$$

x_1	x_2	$h_{\Theta}(x)$
0	0	$g(-30) \approx 0$
0	1	$g(-10) \approx 0$
1	0	$g(-10) \approx 0$
1	1	$g(10) \approx 1$

We can combine AND, OR, NOT, (NOT x_1) AND (NOT x_2) to get XNOR, XOR, etc.

8.6 Multi-Class Classification

Extension of One-vs-All Method!

Example: Say 4 classes, so 4 output units. So, $y \in \mathbb{R}^4$ with 1 in corresponding class and 0 elsewhere.

¹Work out more examples (OR, NOT, (NOT x_1) AND (NOT x_2))!
Check Lecture8.pdf for more details.

Chapter 9

Back Propagation

9.1 Notations

L = total no. of layers in the network

s_l = no. of units (not counting bias unit) in layer l

Binary Classification

$$y = 0 \text{ or } 1$$

1 output unit

$$h_{\Theta}(x) \in \mathbb{R}$$

$$s_L = 1$$

$$K = 1 \quad (\text{for simplification})$$

Multi-class Classification (K classes)

$$y \in \mathbb{R}^K$$

Example: $K = 4$

$$y \in \left\{ \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix} \right\}$$

K output units

$$\begin{aligned} h_{\Theta}(x) &\in \mathbb{R}^K \\ s_L &= K \\ K &\geq 3 \end{aligned}$$

9.2 Cost Function

Logistic Regression

$$J(\theta) = - \left[\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Neural Network

$$h_{\Theta}(x) \in \mathbb{R}^K \quad (h_{\Theta}(x))_i = i^{th} \text{output}$$

$$\begin{aligned} J(\Theta) = & -\frac{1}{m} \left[\sum_{i=1}^m \sum_{k=1}^K \left(y_k^{(i)} \log (h_{\Theta}(x^{(i)}))_k + (1 - y_k^{(i)}) \log (1 - (h_{\Theta}(x^{(i)}))_k) \right) \right] \\ & + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} \left(\Theta_{ji}^{(l)} \right)^2 \end{aligned}$$

9.3 Backpropagation Algorithm

9.3.1 Gradient Computation

$$\begin{aligned} J(\Theta) = & -\frac{1}{m} \left[\sum_{i=1}^m \sum_{k=1}^K \left(y_k^{(i)} \log (h_{\Theta}(x^{(i)}))_k + (1 - y_k^{(i)}) \log (1 - (h_{\Theta}(x^{(i)}))_k) \right) \right] \\ & + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} \left(\Theta_{ji}^{(l)} \right)^2 \\ \min_{\Theta} J(\Theta) \end{aligned}$$

Need code to compute:

- $J(\Theta)$
- $\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$

Given one training example (x, y) :

Forward Propagation:

$$\begin{aligned}
 a^{(1)} &= x \\
 z^{(2)} &= \Theta^{(1)} a^{(1)} \\
 a^{(2)} &= g(z^{(2)}) && (\text{add } a_0^{(2)}) \\
 z^{(3)} &= \Theta^{(2)} a^{(2)} \\
 a^{(3)} &= g(z^{(3)}) && (\text{add } a_0^{(3)}) \\
 z^{(4)} &= \Theta^{(3)} a^{(3)} \\
 h_{\Theta}(x) &= a^{(4)} = g(z^{(4)})
 \end{aligned}$$

Gradient Computation: Back Propagation

Intuition: $\delta_j^{(l)}$ = “error” of node j in layer l

For each output unit ($L = 4$)

$$\delta_j^{(4)} = a_j^{(4)} - y_j$$

Vectorize:

$$\begin{aligned}
 \delta^{(4)} &= a^{(4)} - y \\
 \delta^{(3)} &= (\Theta^{(3)})^T \delta^{(4)} \cdot * g'(z^{(3)}) \\
 g'(z^{(3)}) &= a^{(3)} \cdot * (1 - a^{(3)}) \\
 \delta^{(2)} &= (\Theta^{(2)})^T \delta^{(3)} \cdot * g'(z^{(2)}) \\
 g'(z^{(2)}) &= a^{(2)} \cdot * (1 - a^{(2)})
 \end{aligned}$$

No $\delta^{(1)}$

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = a_j^{(l)} \delta_i^{(l+1)} \quad (\text{ignoring } \lambda; \text{ if } \lambda = 0)$$

Backpropagation Algorithm

Training set $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$

Set $\Delta_{ij}^{(l)} = 0$ (for all l, i, j) \rightarrow accumulators to compute $\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$

For $i = 1$ to m

Set $a^{(1)} = x^{(1)}$

Perform forward propagation to compute $a^{(l)}$ for $l = 2, 3, \dots, L$

Using $y^{(i)}$, compute $\delta^{(L)} = a^{(L)} - y^{(i)}$

Compute $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$

$\Delta_{ij}^{(l)} := \Delta_{ij}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$

$\Delta^{(l)} := \Delta^{(l)} + \delta^{(l+1)} (a^{(l)})^T$ (Vectorized)

$$D_{ij}^{(l)} := \begin{cases} \frac{1}{m} \Delta_{ij}^{(l)} + \lambda \Theta_{ij}^{(l)} & \text{if } j \neq 0 \\ \frac{1}{m} \Delta_{ij}^{(l)} & \text{if } j = 0 \end{cases}$$

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = D_{ij}^{(l)}$$

9.4 Implementation Note

9.4.1 Unrolling Parameters

Advanced optimization functions like ‘fminunc’ assume that inputs ‘theta’, ‘gradient’, etc. are vectors. These are matrices in Neural Networks, so we *unroll* them into vectors.

Example: Neural Network with 3 layers (1 input, 1 hidden, 1 output) and $s_1 = 10$, $s_2 = 10$, $s_3 = 1$.

$$\begin{aligned} \Theta^{(1)} &\in \mathbb{R}^{10 \times 11}, \quad \Theta^{(2)} \in \mathbb{R}^{10 \times 11}, \quad \Theta^{(3)} \in \mathbb{R}^{(1 \times 11)} \\ D^{(1)} &\in \mathbb{R}^{10 \times 11}, \quad D^{(2)} \in \mathbb{R}^{10 \times 11}, \quad D^{(3)} \in \mathbb{R}^{(1 \times 11)} \end{aligned}$$

To unroll:

```
thetaVec = [Theta1(:); Theta2(:); Theta3(:)];  
DVec = [D1(:); D2(:); D3(:)];
```

To go back to matrices:

```
Theta1 = reshape(thetaVec(1:110), 10, 11);  
Theta2 = reshape(thetaVec(111:220), 10, 11);  
Theta3 = reshape(thetaVec(221:231), 1, 11);
```

9.4.2 Learning Algorithm

Have initial parameters $\Theta^{(1)}, \Theta^{(2)}, \Theta^{(3)}$

Unroll to get `initialTheta` to pass to

```
fminunc(@costFunction, initialTheta, options)
```

```
function [jVal, gradientVec] = costFunction(thetaVec)
```

From `thetaVec`, get $\Theta^{(1)}, \Theta^{(2)}, \Theta^{(3)}$ using `reshape`

Use forward/back propagation to compute $D^{(1)}, D^{(2)}, D^{(3)}$ and $J(\Theta)$

Unroll $D^{(1)}, D^{(2)}, D^{(3)}$ to get `gradientVec`

9.5 Gradient Checking

An unfortunate problem is that it might seem like it's working even when there's a bug in the backpropagation algorithm (implementation). Subtle bugs can hence cause higher level of error and perform worse than a bug-free implementation. Here, we try to make sure that our implementation is 100% correct.

9.5.1 Numerical Estimation of gradients

To estimate derivative of $J(\theta)$, we calculate $J(\theta + \epsilon)$ and $J(\theta - \epsilon)$ and take the slope of the line connecting these two points.

$$J'(\theta) \approx \frac{J(\theta + \epsilon) - J(\theta - \epsilon)}{2\epsilon}$$

Use pretty small ϵ ($\approx 10^{-4}$).

The above formula is called the two sided difference. Another formula called the one sided difference is

$$J'(\theta) \approx \frac{J(\theta + \epsilon) - J(\theta)}{\epsilon}$$

But the two sided difference usually gives a slightly more accurate approximation.

Implement:

$$\text{gradApprox} = (J(\text{theta} + \text{EPSILON}) - J(\text{theta} - \text{EPSILON})) / (2 * \text{EPSILON})$$

Parameter *vector* θ

$$\begin{aligned} \theta &\in \mathbb{R}^n \\ \theta &= \begin{bmatrix} \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \\ \frac{\partial}{\partial \theta_1} J(\theta) &\approx \frac{J(\theta_1 + \epsilon, \theta_2, \theta_3, \dots, \theta_n) - J(\theta_1 - \epsilon, \theta_2, \theta_3, \dots, \theta_n)}{2\epsilon} \\ \frac{\partial}{\partial \theta_2} J(\theta) &\approx \frac{J(\theta_1, \theta_2 + \epsilon, \theta_3, \dots, \theta_n) - J(\theta_1, \theta_2 - \epsilon, \theta_3, \dots, \theta_n)}{2\epsilon} \\ &\vdots \\ \frac{\partial}{\partial \theta_n} J(\theta) &\approx \frac{J(\theta_1, \theta_2, \theta_3, \dots, \theta_n + \epsilon) - J(\theta_1, \theta_2, \theta_3, \dots, \theta_n - \epsilon)}{2\epsilon} \end{aligned}$$

Implement:

```

for i = 1:n
    thetaPlus = theta;
    thetaPlus(i) = thetaPlus(i) + EPSILON;
    thetaMinus = theta;
    thetaMinus(i) = thetaMinus(i) - EPSILON;
    gradApprox(i) = (J(thetaPlus) - J(thetaMinus))/(2*EPSILON);
end;

```

Check that $\text{gradApprox} \approx \text{DVec}$ (from backpropagation).

9.5.2 Implementation Note:

- Implement backpropagation to compute `Dvec` (unrolled $D^{(1)}, D^{(2)}, \dots$)
- Implement numerical gradient check to compute `gradApprox`
- Make sure they give similar values
- Turn off gradient checking (much slower) and use backpropagation for learning (much faster)

Important:

- Be sure to disable your gradient checking code before training your classifier. If you run numerical gradient computation on every iteration of gradient descent (or in the inner loop of `costFunction(...)`) your code will be very slow.

9.6 Random Initialization

Initial value of Θ

For gradient descent and advanced optimization methods, we need initial value of Θ .

Consider gradient descent:

Set `initialTheta = zeros(n, 1)` ?

Works fine for linear/logistic regression, but not when we're training a neural network. Why?

$$\begin{aligned}\Theta_{ij}^{(l)} &= 0 \text{ for all } i, j, l \\ a_1^{(2)} &= a_2^{(2)} \\ \delta_1^{(2)} &= \delta_2^{(2)} \\ \frac{\partial}{\partial \Theta_{01}^{(1)}} J(\Theta) &= \frac{\partial}{\partial \Theta_{02}^{(1)}} J(\Theta)\end{aligned}$$

Meaning, after each gradient descent update, parameters corresponding to inputs going into each unit of the next layer are identical. Thus, all the units of the next layer are still computing the same function of the input. The values may change, but will always be the same, preventing the neural network from learning something interesting since it is extremely redundant.

Random Initialization: Symmetry Breaking

To get around this (*symmetric weights*), we do random initialization. Initialize each $\Theta_{ij}^{(l)}$ to a random value in $[-\epsilon, \epsilon]$ (i.e. $-\epsilon \leq \Theta_{ij}^{(l)} \leq \epsilon$).

Example:

```
Theta1 = rand(10,11)*(2*INIT_EPSILON) - INIT_EPSILON
Theta2 = rand(1, 11)*(2*INIT_EPSILON) - INIT_EPSILON
```

A good choice of ϵ is:

$$\epsilon = \frac{\sqrt{6}}{\sqrt{L_{in} + L_{out}}}$$

Where L_{in} and L_{out} are number of units in the layers adjacent to $\Theta^{(l)}$.

9.7 Putting it all together

Pick a network architecture (connectivity pattern between neurons), like number of layers, number of units in each layer, etc. For example, $[3 \rightarrow 5 \rightarrow 4]$ vs $[3 \rightarrow 5 \rightarrow 5 \rightarrow 4]$ vs $[3 \rightarrow 5 \rightarrow 5 \rightarrow 5 \rightarrow 4]$.

No. of input units : Dimension of features $x^{(i)}$

No. of output units : Number of classes

Remember:

$y \in \{1, 2, \dots, 10\}$ becomes

$$y \in \left\{ \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}, \dots, \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix} \right\}$$

For Hidden Layers?

Reasonable Default: 1 hidden layer, or if > 1 hidden layer, have same no. of hidden units in every layer (usually, the more the better, just that it may get computationally expensive). Usually, comparable to the number of features.

Training a neural network:

1. Randomly initialize weights
2. Implement forward propagation to get $h_{\Theta}(x)$ for any x
3. Implement code to compute cost function $J(\Theta)$
4. Implement backpropagation to compute partial derivatives $\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$

```
for i = 1:m
    Perform forward propagation ...
    and backpropagation ...
    using example  $(x^{(i)}, y^{(i)})$ 

    Get activations  $a^{(l)}$  ...
    and delta terms  $\delta^{(l)}$  ...
    for  $l = 2, \dots, L$ 

         $\Delta^{(l)} := \Delta^{(l)} + \delta^{(l+1)}(a^{(l)})^T$ 
    end
end
compute  $\frac{\partial}{\partial \Theta_{jk}^{(l)}} J(\Theta)$ 
```

5. (a) Use gradient checking to compare $J'(\Theta)$ computed using backpropagation vs using numerical estimate.
(b) Then disable gradient checking code.
6. Use gradient descent or advanced optimization method with backpropagation to try to minimize $J(\Theta)$

$J(\Theta)$ in general for a neural network is non-convex, and hence susceptible to local optima. However, it turns out, in practice, this is not usually a huge problem.

Check Lecture9.pdf for more details.

Chapter 10

Applying Machine Learning

10.1 Deciding what to try next

10.1.1 Debugging a Learning Algorithm

Suppose you have implemented regularized linear regression to predict housing prices.

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

However, it makes *unacceptably* large errors in prediction. What to try next?

- Get more training examples (Sometimes doesn't actually help)
- Try smaller set of features (Overfitting?)
- Try additional features (Underfitting?) (Huge project)
- Try adding polynomial features
- Increasing/Decreasing λ

Need to understand what would help.

10.1.2 Machine Learning Diagnostics

Diagnostic: A test that you can run to gain insight at what is/isn't working with the algorithm, and gain guidance as to how best to improve its performance.

Diagnostics can take time to implement, but doing so can be a very good use of your time as it can save huge time later.

10.2 Evaluating your Hypothesis

Just *training error* not a good indicator (Overfitting).

Solution? Plot Hypothesis function? Not feasible with high number of features.

10.2.1 The Standard Way

Split the dataset into two portions - *Training Set* and *Test Set* with about 70%-30% ratio (*randomly* if it is in some order).

Notation:

Training Examples : (x, y)

Test Examples : (x_{test}, y_{test})

No. of Training Examples : m

No. of Test Examples : m_{test}

10.2.2 Train/Test Procedure

1. Learn θ from training data (minimize $J(\theta)$)
2. Compute test set error ($J_{test}(\theta)$) on test examples
 - In case of classification problem, we can also use misclassification error (Also called 0/1 misclassification error)
 - $$\text{err}(h_{\theta}(x), y) = \begin{cases} 1 & \text{if } h_{\theta}(x) \geq 0.5, y == 0 \\ & \text{or if } h_{\theta}(x) < 0.5, y == 1 \\ 0 & \text{otherwise} \end{cases}$$
 - $$\text{test error} = \frac{1}{m_{test}} \sum_{i=1}^{m_{test}} \text{err}(h_{\theta}(x_{test}^{(i)}), y_{test}^{(i)})$$
 - This gives us the proportion of the test data that was misclassified

10.3 Model Selection

Once parameters θ were fit to some training set, the error $J(\theta)$ measured on that data is likely to be lower than the actual generalization error.

Model Selection: Let's say we try to choose what degree polynomial to fit to data (linear, quadratic, cubic, ...)

Notation:

d = Degree of polynomial

For different d , we can have the following:

1. $h_{\theta}(x) = \theta_0 + \theta_1 x$
2. $h_{\theta}(x) = \theta_0 + \theta_1 x + \theta_2 x^2$
3. $h_{\theta}(x) = \theta_0 + \theta_1 x + \dots + \theta_3 x^3$
- \vdots
10. $h_{\theta}(x) = \theta_0 + \theta_1 x + \dots + \theta_{10} x^{10}$

Choose one model out of the above. We can train them separately, and get different $\theta^{(k)}$, and calculate $J_{test}(\theta)$ for each of them, and choose the one with the least test error.

How well does this model generalize? Report $J_{test}(\theta)$.

Problem: $J_{test}(\theta)$ is likely an optimistic estimate of generalization error (our extra parameter d is fit to *test set*).

10.3.1 Evaluating your hypothesis

Divide the data set in three pieces, *training set* (60%), *cross-validation set* (20%), *test set* (20%).

Notations:

Training Examples :	(x, y)
Cross-Validation Examples :	(x_{cv}, y_{cv})
Test Examples :	(x_{test}, y_{test})
No. of Training Examples :	m
No. of Cross-Validation Examples :	m_{cv}
No. of Test Examples :	m_{test}
Training Error :	$J(\theta)$ or $J_{train}(\theta)$
Cross Validation Error :	$J_{cv}(\theta)$
Test Error :	$J_{test}(\theta)$

Use *Cross-Validation Set* to select the model. Estimate generalization error for test set.

10.4 Bias vs Variance

Plot *error* against *degree* of the polynomial for both, *training* and *validation* data set. Usually, training error tends to decrease. Cross-Validation error (or test error) on the other hand, tends to decrease, and then increase again.

To recognise bias/variance problem:

Bias (Underfit) :	J_{train} will be high
	$J_{cv} \approx J_{train}$
Variance (Overfit) :	J_{train} will be low
	$J_{cv} \gg J_{train}$

10.4.1 Effect of regularization

Large λ can cause high bias (underfit), and at very small λ causes high variance (overfit). How to automatically choose a good value of λ .

Choosing λ :

$$\begin{aligned}J(\theta) &= \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) + \frac{\lambda}{2m} \sum_{j=1}^m \theta_j^2 \\J_{train}(\theta) &= \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \\J_{cv}(\theta) &= \frac{1}{2m_{cv}} \sum_{i=1}^{m_{cv}} (h_{\theta}(x_{cv}^{(i)}) - y_{cv}^{(i)}) \\J_{test}(\theta) &= \frac{1}{2m_{test}} \sum_{i=1}^{m_{test}} (h_{\theta}(x_{test}^{(i)}) - y_{test}^{(i)})\end{aligned}$$

Have a model $(h_{\theta}(x), J(\theta))$, have a bunch of values to try for λ like 0, 0.01, 0.02, 0.04, 0.08, \dots , 10 (roughly doubling). Train (minimize $J(\theta)$) for all values. Use $J_{cv}(\theta)$ to choose a λ . Report generalization error as $J_{test}(\theta)$.

Bias/Variance as a function of λ

Plot $J_{train}(\theta)$ and $J_{cv}(\theta)$ as a function of λ . Usually, $J_{train}(\theta)$ tends to increase while $J_{cv}(\theta)$ tends to decrease at first but increase again later.

10.5 Learning Curves

Useful thing to plot, to sanity check, to improve the performance, to diagnose for bias/variance, \dots

Plot $J_{train}(\theta)$, and $J_{cv}(\theta)$ as a function of m , the number of training examples. Measure $J_{train}(\theta)$ only on data used to train!. Measure $J_{cv}(\theta)$ on all cross validation data!

Usually, average training error will increase with m . Usually, average validation error will decrease with m .

10.5.1 High Bias

Usually: Validation error will decrease at first but then will saturate/flatten out soon. Training error will increase at first but then will saturate/flatten out soon, close to validation error.

Note: If a learning algorithm is suffering from high bias, getting more training data will not (by itself) help much.

10.5.2 High Variance

Usually: Training error will increase but still be pretty low. Validation error will decrease but still be pretty high, far from training error.

Note: If a learning algorithm is suffering from high variance, getting more training examples is indeed likely to help.

10.6 Deciding what to try next (Revisited)

Get more training examples	fixes high variance
Try smaller set of features	fixes high variance
Try additional features	fixes high bias (not always)
Try adding polynomial features	fixes high bias (not always)
Increasing/Decreasing λ	fixes bias/variance

10.6.1 Neural Networks

Small neural network, computationally cheaper, but possibly underfitting. *Large* neural networks, computationally more expensive, and possibly overfitting. Use regularization to address overfitting.

Often, large neural network with regularization works better (but may be computationally expensive).

Large network can comprise of a single large hidden layer, or two moderate hidden layers, or three average hidden layers, etc.

Check Lecture10.pdf for more details.

Chapter 11

Machine Learning System Design

11.1 Prioritizing what to work on: Spam Classification Example

Supervised Learning! x = features of email. y = spam(1) or not spam(0).

Features x : Choose 100 words indicative of spam/not spam.

Example: deal, buy, discount, andrew, now, ...

Feature vector:

$$x = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \\ \vdots \\ 1 \end{bmatrix} \begin{bmatrix} \textit{andrew} \\ \textit{buy} \\ \textit{deal} \\ \textit{discount} \\ \vdots \\ \textit{now} \end{bmatrix}$$

0/1 depending on whether a 'word' appears or not in the email.

$$x_j = \begin{cases} 1 & \text{if word } j \text{ appears in email} \\ 0 & \text{otherwise} \end{cases}$$

Note: In practice, take most frequently occurring $n(10,000 \text{ to } 50,000)$ words in training set, instead of manually picking 100 words.

11.1.1 Time Management?

- Collect lots of data
 - E.g. *honeypot* project.
- Develop sophisticated features based on email routing information (from email header).
- Develop sophisticated features for message body, e.g. should *discount* and *discounts* be treated as same word? How about *deal* and *Dealer*? Features about punctuation?
- Develop sophisticated algorithm to detect misspellings (e.g. m0rtgage, med1cine, w4tches.)

11.2 Error Analysis

11.2.1 Recommended approach

- Start with a simple and quick algorithm, implement and test it on cross-validation data.
- Plot learning curves to decide if more data, more features, etc.
- Error Analysis: Manually examine the examples (in cross validation set) that your algorithm made errors on. See if you spot any systematic trend in what type of examples it is making errors on.

$m_{cv} = 500$ examples in cross validation set, algorithm misclassifies 100 emails. Manually examine 100 errors, and categorized them based on:

1. What type of email it is (pharma, replica, phishing, ...)
2. What features would have helped (Deliberate misspellings, Unusual email routing, Unusual punctuation)

11.2.2 Numerical Evaluation

Have a way of *numerical evaluation* like accuracy/error, to tell how the learning algorithm is doing.

Example: Should discount/discounts/discounting be treated as the same word? Can use *stemming* software (E.g. "Porter Stemmer"). It can hurt because the software can mistake universe/university to be the same.

So, error analysis may not be helpful for deciding if this is likely to improve performance. Only solution is to try and see if it works.

Need numerical evaluation (e.g. cross validation error) of algorithm's performance with and without stemming.

11.3 Skewed Classes

Cancer classification example: Train logistic regression model $h_{\theta}(x)$ with 1% error on test set. Good? What if only 0.50% patients in training/test set have cancer? Not so impressive anymore! Even predicting *No Cancer* all the time will get 0.50% error. This is a case of *skewed* classes.

Going from 99.2% to 99.5% accuracy is actually doing something useful? Can't say in case of skewed classes.

11.3.1 Precision/Recall

$y = 1$ in presence of rare class that we want to detect.

Notations

Notations	Actual Class	Predicted Class
True Positive	1	1
True Negative	0	0
False Positive	0	1
False Negative	1	0

Precision

Out of all the predicted *positives*, what fraction is *actually* positive.

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

Recall

Out of all the actual *positives*, what fraction is *predicted* positive.

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

Now, if we predict no cancer always, then $\text{recall} = 0$. These metrics usually help us evaluate our algorithm even in case of really skewed classes.

11.3.2 Trade-Off

Consider same example of cancer, we train logistic regression $h_\theta(x)$ which gives the probability of having cancer.

Suppose we want to predict cancer ($y = 1$) only if we're very sure. One way to do this, is predict $y = 1$ when $h_\theta(x) \geq 0.7$ instead of 0.5 and $y = 0$ otherwise. This will give us *higher* precision, but *lower* recall.

Suppose we want to avoid missing too many cases of cancer, we can predict $y = 1$ when $h_\theta(x) \geq 0.3$. This will give us *higher* recall, but *lower* precision.

How to choose the threshold?

F_1 Score (F Score)

We have lost our 'single number' evaluation metric, we now have precision and recall. What are our options?

- Average of precision and recall ($\frac{P+R}{2}$)? Not a good idea! (Extremes)
- F_1 Score = $2\frac{PR}{P+R}$ better choice, commonly used

Use F_1 score on cross-validation set to decide the threshold.

11.3.3 Working with Large Data

Consider cases where more data will be helpful.

Say, feature $x \in \mathbb{R}^{n+1}$ has sufficient information to predict y accurately. (Counter-examples: predict housing price from only size or area).

A useful test is, given x , can a human expert confidently predict y ?

Alongside, Suppose we use a learning algorithm with a large number of parameters (low bias algorithms).

Using a very large training set, it is unlikely to overfit.

Check Lecture11.pdf for more details.

Chapter 12

Support Vector Machines

Another popular supervised learning algorithm.

12.1 Optimization Objective

12.1.1 Alternative view of logistic regression

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

If $y = 1$, we want $h_{\theta}(x) \approx 1$, $\theta^T x \gg 0$

If $y = 0$, we want $h_{\theta}(x) \approx 0$, $\theta^T x \ll 0$

Cost of example:

Each example (x, y) contributes:

$$\begin{aligned} & - (y \log(h_{\theta}(x)) + (1 - y) \log(1 - h_{\theta}(x))) \\ &= -y \log\left(\frac{1}{1 + e^{\theta^T x}}\right) - (1 - y) \log\left(1 - \frac{1}{1 + e^{\theta^T x}}\right) \end{aligned}$$

If $y = 1$ (want $\theta^T x \gg 0$): we get $-\log\left(\frac{1}{1 + e^{-z}}\right)$ where $z = \theta^T x$. For *SVM* modify cost function to be flat 0 for $x > 1$ and a straight line to the left (with negative slope). Denote by $cost_1(z)$.

If $y = 0$ (want $\theta^T x \ll 0$): we get $-\log \left(1 - \frac{1}{1 + e^{-z}} \right)$ where $z = \theta^T x$.
 For *SVM* modify cost function to be flat 0 for $x < -1$ and a straight line to the left (with positive slope). Denote by $cost_0(z)$.

Logistic Regression:

$$\min_{\theta} \frac{1}{m} \left[\sum_{i=1}^m (y^{(i)} (-\log(h_{\theta}(x^{(i)}))) + (1 - y^{(i)}) (-\log(1 - h_{\theta}(x^{(i)})))) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Support Vector Machine:

Modify costs first

$$\min_{\theta} \frac{1}{m} \left[\sum_{i=1}^m (y^{(i)} cost_1(\theta^T x^{(i)}) + (1 - y^{(i)}) cost_0(\theta^T x^{(i)})) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Get rid of $\frac{1}{m}$ term. Only a constant, same optimal value!

$$\min_{\theta} \left[\sum_{i=1}^m (y^{(i)} cost_1(\theta^T x^{(i)}) + (1 - y^{(i)}) cost_0(\theta^T x^{(i)})) \right] + \frac{\lambda}{2} \sum_{j=1}^n \theta_j^2$$

By convention, instead of optimizing $A + \lambda B$, optimize $CA + B$ where C is somewhat equivalent of $\frac{1}{\lambda}$, again, same optimal value!

$$\min_{\theta} C \left[\sum_{i=1}^m (y^{(i)} cost_1(\theta^T x^{(i)}) + (1 - y^{(i)}) cost_0(\theta^T x^{(i)})) \right] + \frac{1}{2} \sum_{j=1}^n \theta_j^2$$

12.1.2 SVM Hypothesis

$$\min_{\theta} C \left[\sum_{i=1}^m (y^{(i)} cost_1(\theta^T x^{(i)}) + (1 - y^{(i)}) cost_0(\theta^T x^{(i)})) \right] + \frac{1}{2} \sum_{j=1}^n \theta_j^2$$

Hypothesis:

Unlike logistic regression, SVM does not give probability. It just makes prediction.

$$h_{\theta}(x) = \begin{cases} 1 & \text{if } \theta^T x \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

12.2 Large Margin

12.2.1 Intuition

If $y = 1$, we want $\theta^T x \geq 1$ (not just ≥ 0)

If $y = 0$, we want $\theta^T x \leq -1$ (not just < 0)

This builds safety margin factor in between.

SVM Decision Boundary

If $C \gg 1$, say 100,000, we will get

$$\min \left(C \times 0 + \frac{1}{2} \sum_{j=1}^n \theta_j^2 \right) \quad \text{such that} \quad \begin{cases} \theta^T x^{(i)} \geq 1 & \text{if } y^{(i)} = 1 \\ \theta^T x^{(i)} \leq -1 & \text{if } y^{(i)} = 0 \end{cases} \quad \text{Very interesting decision boundary!}$$

Linearly Separable Case: It will choose the decision boundary such that it maximises margin of the SVM.

Outliers: If C is indeed very large it will be affected by outliers easily but if C is not too large, it will still give a reasonable boundary.

12.2.2 Math

Vector Inner Product

$u^T v$ (or dot product)

$$u = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix}$$

$$v = \begin{bmatrix} v_1 \\ v_2 \end{bmatrix}$$

$\|u\|$ = length of vector u

$$= \sqrt{u_1^2 + u_2^2} \in \mathbb{R}$$

p = projection of vector v on vector u

$$u^T v = p \|u\|$$

$$u^T v = u_1 v_1 + u_2 v_2$$

$$u^T v = v^T u$$

SVM Decision Boundary

$$\min_{\theta} \frac{1}{2} \sum_{j=1}^n \theta_j^2 \text{ s.t. } \begin{cases} \theta^T x^{(i)} \geq 1 & \text{if } y^{(i)} = 1 \\ \theta^T x^{(i)} \leq -1 & \text{if } y^{(i)} = 0 \end{cases}$$

Simplifications: Only for the purpose of understanding here!

$$\theta_0 = 0$$

$$n = 2$$

This gives, $\min_{\theta} \frac{1}{2} (\theta_1^2 + \theta_2^2) = \min_{\theta} \frac{1}{2} \left(\sqrt{\theta_1^2 + \theta_2^2} \right)^2 = \min_{\theta} \frac{1}{2} \|\theta\|^2$
 $\theta^T x^{(i)} = ?$

$$\begin{aligned} p^{(i)} &= \text{projection of vector } x^{(i)} \text{ on vector } \theta \\ \theta^T x^{(i)} &= p^{(i)} \|\theta\| \\ &= \theta_1 x_1^{(i)} + \theta_2 x_2^{(i)} \end{aligned}$$

Optimization Objective: Modified as follows:

$$\min_{\theta} \frac{1}{2} \sum_{j=1}^n \theta_j^2 = \frac{1}{2} \|\theta\|^2 \text{ s.t. } \begin{cases} p^{(i)} \|\theta\| \geq 1 & \text{if } y^{(i)} = 1 \\ p^{(i)} \|\theta\| \leq -1 & \text{if } y^{(i)} = 0 \end{cases}$$

where $p^{(i)}$ is the projection of $x^{(i)}$ on vector θ . $p^{(i)}$ will be bigger with large margin helping minimize the $\frac{1}{2} \|\theta\|^2$.

12.3 Kernels

12.3.1 Non-linear Decision Boundary

Predict $y = 1$ if $\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots \geq 0$.

New Notation: $h\theta = \theta_0 + \theta_1 f_1 + \theta_2 f_2 + \dots$ where possibly, $f_1 = x_1$, $f_2 = x_2$, $f_3 = x_1 x_2$, $f_4 = x_1^2, \dots$

Is there a better/different choice of features f_1, f_2, f_3, \dots ?

Given x , choose 3 random landmarks l s and define new features as:

$$\begin{aligned} f_1 = \text{similarity}(x, l^{(1)}) &= \exp\left(\frac{-\|x - l^{(1)}\|^2}{2\sigma^2}\right) \\ f_2 = \text{similarity}(x, l^{(2)}) &= \exp\left(\frac{-\|x - l^{(2)}\|^2}{2\sigma^2}\right) \\ f_3 = \text{similarity}(x, l^{(3)}) &= \exp\left(\frac{-\|x - l^{(3)}\|^2}{2\sigma^2}\right) \end{aligned}$$

Similarity function is *Kernel* functions. This particular one is *Gaussian* Kernel.

Kernels and Similarity

If $x \approx l^{(1)}$: $f_1 \approx 1$

If x is far from $l^{(1)}$: $f_1 \approx 0$

Measures closeness to landmarks.

Setting σ^2 low, the value of f_1 falls to 0 very rapidly, and vice-versa.

Predict 1 when $\theta_0 + \theta_1 f_1 + \theta_2 f_2 + \theta_3 f_3 \geq 0$ Say $\theta_0 = -0.5$, $\theta_1 = 1$, $\theta_2 = 1$, $\theta_3 = 0$, then points close to $l^{(1)}$ or $l^{(2)}$ will give 1 creating complex non-linear boundary.

12.3.2 Choosing the Landmarks

Where to get landmarks? How many?

One choice would be one landmark per location of each feature. Given m examples, we have m landmarks.

Given an example x , compute f_1, f_2, \dots using kernels and get the feature vector f (include $f_0 = 1$). For training examples, at least one feature would be 1.

Hypothesis: Given x , compute $f \in \mathbb{R}^{m+1}$

Predict 'y=1' if $\theta^T f \geq 0$

Predict 'y=0' otherwise

Training:

$$\min_{\theta} C \left[\sum_{i=1}^m (y^{(i)} \text{cost}_1(\theta^T f^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T f^{(i)})) \right] + \frac{1}{2} \sum_{j=1}^n \theta_j^2$$

$n = m$

Slight modification in last term! Instead of using $\theta^T \theta$, use $\theta^T M \theta$ where M is a matrix depending on the kernel done for computational expense.

Can use Kernels in other algorithms but tricks (such as M) don't go well with other algorithms so it can be slow/expensive.

12.3.3 SVM Parameters:

$$C \left(= \frac{1}{\lambda} \right)$$

Large C : Lower Bias, High Variance

Small C : Higher Bias, Low Variance

$$\sigma^2$$

Large σ^2 : Features f_i vary more smoothly
: Higher bias, Low Variance

Small σ^2 : Features f_i vary less smoothly
: Lower Bias, High Variance

12.4 Using SVM

Use SVM software package (e.g. liblinear, libsvm, ...) to solve for parameters θ .

Need to specify:

- Choice of parameter C
- Choice of kernel (similarity function):
 - Linear Kernel is just another name of no kernel.
 - Gaussian Kernel. Need to choose σ^2 .

```

function f = kernel(x1, x2)
    f = exp  $\left(-\frac{\|x1 - x2\|^2}{2\sigma^2}\right)$ 
return

```

Note: Do perform feature scaling before using the Gaussian Kernel.

- Other choices of kernel

Note: Not all similarity functions make valid kernels. (Need to satisfy a technical condition called *Mercer's Theorem* to make sure SVM packages' optimizations run correctly, and do not diverge)

Many off-the-shelf kernels available:

- * Polynomial Kernel: $k(x, l) = (x^T l + \text{constant})^{\text{degree}}$ (almost always performs worse than Gaussian Kernel, and not used that much and used only where x and l are all strictly positive)
- * More esoteric: String kernel, chi-square kernel, histogram intersection kernel, ...

12.4.1 Multi-Class Classification

Many SVM packages already have built-in multi-class classification functionality. Otherwise, use one vs all method.

12.4.2 Logistic Regression vs SVM

If $n \geq m$ (say $n = 10,000$ and $m = 10$ to 1000), use logistic regression or SVM without kernel (linear kernel).

If n is small (say upto 1000), m is intermediate (say upto $10,000$), use SVM with Gaussian Kernel.

If n is small (say upto 1000), m is large (say $50,000+$), SVM with Gaussian Kernels may start struggling. Create/Add more features, then use logistic regression or SVM without a kernel.

Neural Network likely to work well for most of these settings, but may be slower to train.

SVM is a convex optimization problem, so no need to worry about local optima.

Check Lecture12.pdf for more details.