

# Notes

## Machine Learning by Andrew Ng on Coursera

Sparsh Jain

September 25, 2020

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Supervised Learning . . . . .	3
1.1.1	Regression Problem . . . . .	3
1.1.2	Classification Problems . . . . .	3
1.2	Unsupervised Learning . . . . .	3
<b>2</b>	<b>Linear Regression with One Variable</b>	<b>4</b>
2.1	Notations . . . . .	4
2.2	Supervised Learning . . . . .	4
2.3	Gradient Descent . . . . .	5
2.4	Gradient Descent for Linear Regression . . . . .	5
<b>3</b>	<b>Linear Algebra</b>	<b>6</b>
3.1	Matrix . . . . .	6
3.2	Vector . . . . .	6
3.3	Addition and Scalar Multiplication . . . . .	7
3.4	Matrix Matrix Multiplication . . . . .	7
3.5	Inverse and Transpose . . . . .	7
<b>4</b>	<b>Linear Regression with Multiple Variables</b>	<b>9</b>
4.1	Notations . . . . .	9
4.2	Hypothesis . . . . .	9
4.3	Gradient Descent . . . . .	10
4.3.1	Feature Scaling . . . . .	10
4.3.2	Mean Normalization . . . . .	11
4.3.3	Learning Rate . . . . .	11
4.4	Features and Polynomial Regression . . . . .	11
4.4.1	Features . . . . .	11
4.4.2	Polynomial Regression . . . . .	11
4.5	Normal Equation . . . . .	12
4.5.1	Non Invertibility of $X^T X$ . . . . .	12

<b>5</b>	<b>Octave Tutorial</b>	<b>13</b>
<b>6</b>	<b>Classification</b>	<b>14</b>
6.1	Logistic Regression . . . . .	14
6.2	Decision Boundary . . . . .	15
6.3	Cost Function . . . . .	15
6.4	Advanced Optimization: . . . . .	17
6.5	Multi-Class Classification . . . . .	18
6.5.1	One vs All . . . . .	18
	<b>Appendices</b>	<b>19</b>

# Chapter 1

## Introduction

*Machine learning* (task, experience, performance) can be classified into *Supervised* and *Unsupervised* learning.

### 1.1 Supervised Learning

Supervised learning can be basically classified into *Regression* and *Classification* problems.

#### 1.1.1 Regression Problem

Regression problems work loosely on continuous range of outputs.

#### 1.1.2 Classification Problems

Classification problems work loosely on discrete range of outputs.

### 1.2 Unsupervised Learning

An example is *Clustering Problem*.

---

Check Lecture1.pdf for more details.

# Chapter 2

## Linear Regression with One Variable

### 2.1 Notations

$m$  = number of training examples

$x$ 's = 'input' variables / features

$y$ 's = 'output' variables / 'target' variables

$(x, y)$  = single training example

$(x^{(i)}, y^{(i)}) = i^{th}$  example

### 2.2 Supervised Learning

We have a data set (*Training Set*).

Training Set  $\rightarrow$  Learning Algorithm  $\rightarrow h$  (*hypothesis*, a function  $X \rightarrow Y$ )

**To Represent  $h$**

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

**Cost**

$$\underset{\theta_0, \theta_1}{\text{minimize}} \frac{1}{2m} \sum_1^m (h_{\theta}(x) - y)^2$$

## Cost Function

Squared Error Cost Function

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_1^m (h_{\theta}(x) - y)^2$$

$$\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$$

## 2.3 Gradient Descent

Finds local optimum:

1. Start with some value
2. Get closer to optimum

### Algorithm

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \quad \forall j$$

where  $\alpha$  = learning rate

### Important!

Simultaneous Update!

$$\begin{aligned} temp_j &:= \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \quad \forall j \\ \theta_j &:= temp_j \quad \forall j \end{aligned}$$

## 2.4 Gradient Descent for Linear Regression

Cost function for linear regression is convex!

*Batch Gradient Descent:* Each step of gradient descent uses all training examples.

---

Check Lecture2.pdf for more details.

# Chapter 3

## Linear Algebra

### 3.1 Matrix

Rectangular array of numbers:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

**Dimension of the matrix:** #rows x #cols (2 x 3)

**Elements of the matrix:**

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

$A_{ij}$  = “ $i, j$  entry” in the  $i^{th}$  row,  $j^{th}$  col

### 3.2 Vector

An  $n \times 1$  matrix.

$$y = \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix}$$

$$y_i = i^{th} \text{ element}$$

**Note:** Uppercase for matrices, lowercase for vectors.

### 3.3 Addition and Scalar Multiplication

Add/Subtract (element by element) matrices of same dimension only!

Multiply/Divide (all elements) a matrix by scalar!

### 3.4 Matrix Matrix Multiplication

$m \times n$  matrix multiplied by  $n \times o$  matrix gives a  $m \times o$  matrix.

#### Properties

1. Matrix Multiplication is *not* Commutative.
2. Matrix Multiplication is Associative.
3. *Identity Matrix (I)*: 1's along diagonal, 0's everywhere else in an  $n \times n$  matrix.  $AI = IA = A$ .

### 3.5 Inverse and Transpose

#### Inverse

Only square ( $n \times n$ ) matrices *may* have an inverse.

$$AA^{-1} = A^{-1}A = I$$

Matrices that don't have an inverse are *singular* or *degenerate* matrices.

#### Transpose

Let  $A$  be an  $m \times n$  matrix and let  $B = A^T$ , then

$$B_{ij} = A_{ji}$$



Example:

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$
$$B = A^T = \begin{bmatrix} 1 & 4 \\ 2 & 5 \\ 3 & 6 \end{bmatrix}$$

---

Check Lecture3.pdf for more details.

# Chapter 4

## Linear Regression with Multiple Variables

### 4.1 Notations

$n$  = number of features

$x^{(i)}$  = input (features) of  $i^{th}$  training example

$x_j^{(i)}$  = value of feature  $j$  of  $i^{th}$  training example

### 4.2 Hypothesis

Previously:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Now:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

For convinience, define  $x_0 = 1$ . So

$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$h_{\theta}(x) = \theta^T x$$

## 4.3 Gradient Descent

Hypothesis :  $h_{\theta}(x) = \theta^T x$

Parameters :  $\theta$

$$= \theta_0 x_0 + \theta_1 x_1 + \cdots + \theta_n x_n$$

$$= \theta_0, \theta_1, \dots, \theta_n$$

Cost Function :  $J(\theta) = J(\theta_0, \theta_1, \dots, \theta_n)$

$$= \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Gradient Descent :

Repeat{

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

$$= \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1, \dots, \theta_n)$$

$$= \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

} (simultaneously update  $\forall j = 0, 1, \dots, n$ )

### 4.3.1 Feature Scaling

**Idea:** Make sure features are on a similar scale.

Get every feature into approximately a  $-1 \leq x_i \leq 1$  range.

### 4.3.2 Mean Normalization

Replace  $x_i$  with  $x_i - \mu_i$  to make features have approximately zero mean (Do not apply to  $x_0 = 1$ ).

#### General Rule

$$x_i \leftarrow \frac{x_i - \mu_i}{S_i}$$

where

$\mu_i$  = average value of  $x_i$

$S_i$  = range (max - min) *or*

=  $\sigma$ (standard deviation)

### 4.3.3 Learning Rate

$J(\theta)$  should decrease after every iteration. #iterations vary a lot.

Example *Automatic Convergence Test*: Declare convergence if  $J(\theta)$  decreases by less than  $\epsilon$  (say  $10^{-3}$ ) in one iteration.

If  $J(\theta)$  increases, use smaller  $\alpha$ . Too small  $\alpha$  means slow convergence.

To choose  $\alpha$ , try ..., 0.001, 0.003, 0.01, 0.03, 0.1, 0.3, 1, ...

## 4.4 Features and Polynomial Regression

### 4.4.1 Features

Get an insight in your problem and choose better features (may even combine/separate features).

Ex: size = length  $\rightarrow$  breadth.

### 4.4.2 Polynomial Regression

Ex:

$$x_1 = size$$

$$x_2 = size^2$$

$$x_3 = size^3$$

## 4.5 Normal Equation

Solve for  $\theta$  analytically!

$$\begin{aligned}x^{(i)} &= \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix} && \in \mathbb{R}^{n+1} \\X &= \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ \vdots \\ (x^{(m)})^T \end{bmatrix} && \in \mathbb{R}^{m \times (n+1)} \\&= \begin{bmatrix} x_0 & x_1 & \dots & x_n \end{bmatrix} && \in \mathbb{R}^{m \times (n+1)} \\y &= \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix} && \in \mathbb{R}^m \\\theta &= (X^T X)^{-1} X^T y\end{aligned}$$

Inverse of a matrix grows as  $O(n^3)$ , use wisely.

### 4.5.1 Non Invertibility of $X^T X$

Use 'pinv' function in Octave (pseudo-inverse) instead of 'inv' function (inverse).

If  $X^T X$  is non-invertible, common causes are

1. Redundant features (linearly dependent)
2. Too many features ( $m \leq n$ ). In this case, delete some features or use *regularization*

---

Check Lecture4.pdf for more details.

# Chapter 5

## Octave Tutorial

Check Lecture5.pdf for more details.

# Chapter 6

## Classification

Classify into categories (binary or multiple).

### 6.1 Logistic Regression

$$\begin{aligned} 0 &\leq h_{\theta}(x) \leq 1 \\ h_{\theta}(x) &= g(\theta^T x) \\ g(z) &= \frac{1}{1 + e^{-z}} \quad g \text{ is called a } \textit{sigmoid} \text{ function or a } \textit{logistic} \text{ function.} \\ h_{\theta}(x) &= \frac{1}{1 + e^{\theta^T x}} \end{aligned}$$

### Interpretation of Hypothesis Output

$h_{\theta}(x)$  = estimated probability that  $y = 1$  on input  $x$

$h_{\theta}(x) = P(y = 1|x; \theta)$  = probability that  $y = 1$ , given  $x$ , parameterized by  $\theta$

$$P(y = 0|x; \theta) + P(y = 1|x; \theta) = 1$$

## 6.2 Decision Boundary

$$\begin{array}{ll} \text{Predict: } y = 1 & \text{if } h_{\theta}(x) \geq 0.5 \\ & (\theta^T x \geq 0) \\ \text{Predict: } y = 0 & \text{if } h_{\theta}(x) < 0.5 \\ & (\theta^T x < 0) \\ \theta^T x = 0 & \text{is the } \textit{decision boundary}. \end{array}$$

## Non-linear Decision Boundaries

Use same technique as polynomial regression for features.

## 6.3 Cost Function

Training Set :  $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$x_0 = 1$$

$$y \in \{0, 1\}$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$



## How to choose parameter $\theta$ ?

### Linear Regression:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

$$\text{Cost}(h_{\theta}(x), y) = \frac{1}{2} (h_{\theta}(x) - y)^2$$

### Logistic Regression:

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & y = 1 \\ -\log(1 - h_{\theta}(x)) & y = 0 \end{cases}$$

$$\text{Cost}(h_{\theta}(x), y) = 0 \text{ if } h_{\theta}(x) = y$$

$$\text{Cost}(h_{\theta}(x), y) \rightarrow \inf \text{ if } y = 0 \text{ and } h_{\theta}(x) \rightarrow 1$$

$$\text{Cost}(h_{\theta}(x), y) \rightarrow \inf \text{ if } y = 1 \text{ and } h_{\theta}(x) \rightarrow 0$$

**Note:**  $y = 0$  or  $1$  always.

$$\text{Cost}(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$$

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[ \sum_{i=1}^m (y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))) \right] \end{aligned}$$

To fit parameters  $\theta$ :

$$\underset{\theta}{\text{minimize}} J(\theta)$$

To make a prediction given a new  $x$ :

$$\text{Output } h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

### Gradient Descent:

Simultaneously update all  $\theta_j$

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Plug in the derivative

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Don't forget feature scaling!

## 6.4 Advanced Optimization:

Something better than gradient descent:

1. Conjugate Gradient
2. BFGS
3. L-BFGS

Advantages:

1. No need to manually pick  $\alpha$
2. Often faster than gradient descent

Disadvantages:

1. More complex

Use libraries! Beware of bad implementations!

**How to use:** We first need to provide a function that evaluates the following two functions for a given input value of  $\theta$ .

1.  $J(\theta)$
2.  $\frac{\partial}{\partial \theta_j} J(\theta)$

```

% We can write a single function that can return both of these:
function [jVal, gradient] = costFunction(theta)
    jVal = [...code to compute J(theta)...];
    gradient = [...code to compute derivative of J(theta)...];
end

```

Then we can use octave's *fminunc()* optimization algorithm along with the *optimset()* function that creates an object containing the options we want to send to *fminunc()*.

```

options = optimset('GradObj', 'on', 'MaxIter', 100);
initialTheta = zeros(2,1);
[optTheta, functionVal, exitFlag] = fminunc(@costFunction,
    initialTheta, options);

```

## 6.5 Multi-Class Classification

### 6.5.1 One vs All

Build a separate binary classifier  $h_{\theta}^{(i)}(x)$  for each class against all other classes.

$$h_{\theta}^{(i)} = P(y = i|x; \theta) \quad \forall i$$

On a new input  $x$ , to make a prediction, pick the class  $i$  that maximizes  $h_{\theta}^{(i)}(x)$

---

Check Lecture6.pdf for more details.

# Appendices