Notes
Machine Learning by Andrew Ng on Coursera

Sparsh Jain

November 23, 2020

Contents

Chapter 1

Introduction

Machine learning (task, experience, performance) can be classified into *Supervised* and *Unsupervised* learning.

1.1 Supervised Learning

Supervised learning can be basically classified into *Regression* and *Classification* problems.

1.1.1 Regression Problem

Regression problems work loosely on continuous range of outputs.

1.1.2 Classification Problems

Classification problems work loosely on discrete range of outputs.

1.2 Unsupervised Learning

An example is *Clustering Problem*.

Check Lecture1.pdf for more details.

Part I

Supervised Learning

Chapter 2

Linear Regression with One Variable

2.1 Notations

$$m = \text{number of training examples}$$
$$x\text{'s} = \text{'input' variables / features}$$
$$y\text{'s} = \text{'output' variables / 'target' variables}$$
$$(x, y) = \text{single training example}$$
$$(x^{(i)}, y^{(i)}) = i^{th} \text{ example}$$

2.2 Supervised Learning

We have a data set (*Training Set*).

Training Set → Learning Algorithm → h (*hypothesis*, a function $X \to Y$)

To Represent h

$$h_\theta(x) = \theta_0 + \theta_1 x$$

Cost

$$\underset{\theta_0, \theta_1}{\text{minimize}} \frac{1}{2m} \sum_1^m (h_\theta(x) - y)^2$$

Cost Function

Squared Error Cost Function

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_1^m (h_\theta(x) - y)^2$$

$$\underset{\theta_0, \theta_1}{\text{minimize}} \, J(\theta_0, \theta_1)$$

2.3 Gradient Descent

Finds local optimum:

1. Start with some value

2. Get closer to optimum

Algorithm

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \; \forall j$$

where α = learning rate

Important!

Simultaneous Update!

$$temp_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \; \forall j$$
$$\theta_j := temp_j \; \forall j$$

2.4 Gradient Descent for Linear Regression

Cost function for linear regression is convex!

Batch Gradient Descent: Each step of gradient descent uses all training examples.

Check Lecture2.pdf for more details.

Chapter 3

Linear Algebra

3.1 Matrix

Rectangular array of numbers:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

Dimension of the matrix: #rows x #cols (2 x 3)

Elements of the matrix:

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

$A_{ij} =$ "i, j entry" in the i^{th} row, j^{th} col

3.2 Vector

An $n \times 1$ matrix.

$$y = \begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix}$$

$y_i = i^{th}$ element

Note: Uppercase for matrices, lowercase for vectors.

3.3   Addition and Scalar Multiplication

Add/Subtract (element by element) matrices of same dimention only!
Multiply/Divide (all elements) a matrix by scalar!

3.4   Matrix Matrix Multiplication

$m \times n$ matrix multiplied by $n \times o$ matrix gives a $m \times o$ matrix.

Properties

1. Matrix Multiplication is *not* Commutative.

2. Matrix Multiplication is Associative.

3. *Identity Matrix (I):* 1's along diagonal, 0's everywhere else in an $n \times n$ matrix. $AI = IA = A$.

3.5   Inverse and Transpose

Inverse

Only square ($n \times n$) matrices *may* have an inverse.

$$AA^{-1} = A^{-1}A = I$$

.

Matrices that don't have an inverse are *singular* or *degenerate* matrices.

Transpose

Let A be an $m \times n$ matrix and let $B = A^T$, then

$$B_{ij} = A_{ji}$$

Example:

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

$$B = A^T = \begin{bmatrix} 1 & 4 \\ 2 & 5 \\ 3 & 6 \end{bmatrix}$$

Check Lecture3.pdf for more details.

Chapter 4

Linear Regression with Multiple Variables

4.1 Notations

$$n = \text{number of features}$$
$$x^{(i)} = \text{input (features) of } i^{th} \text{ training example}$$
$$x_j^{(i)} = \text{value of feature } j \text{ of } i^{th} \text{ training example}$$

4.2 Hypothesis

Previously:

$$h_\theta(x) = \theta_0 + \theta_1 x$$

Now:

$$h_\theta(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

For convinience, define $x_0 = 1$. So

$$h_\theta(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1} \qquad\qquad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$h_\theta(x) = \theta^T x$$

4.3 Gradient Descent

$$
\begin{aligned}
\text{Hypothesis} &: h_\theta(x) = \theta^T x & &= \theta_0 x_0 + \theta_1 x_1 + \cdots + \theta_n x_n \\
\text{Parameters} &: \theta & &= \theta_0, \theta_1, \ldots, \theta_n \\
\text{Cost Function} &: J(\theta) = J(\theta_0, \theta_1, \ldots, \theta_n) & &= \frac{1}{2m} \sum_{i=1}^{m} (h_\theta(x^{(i)}) - y^{(i)})^2
\end{aligned}
$$

Gradient Descent:

```
repeat{
```
$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$
$$= \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1, \ldots, \theta_n)$$
$$= \theta_j - \alpha \frac{1}{m} \sum_{i=1}^{m} (h_\theta(x^{(i)} - y^{(i)}) x_j^{(i)})$$
```
}(simultaneously update ∀j = 0, 1,..., n)
```

4.3.1 Feature Scaling

Idea: Make sure features are on a similar scale.

Get every feature into approximately a $-1 \le x_i \le 1$ range.

4.3.2 Mean Normalization

Replace x_i with $x_i - \mu_i$ to make features have approximately zero mean (Do not apply to $x_0 = 1$).

General Rule

$$x_i \leftarrow \frac{x_i - \mu_i}{S_i}$$

where

$$\mu_i = \text{average value of } x_i$$
$$S_i = \text{range (max - min)} \qquad \qquad \textit{or}$$
$$= \sigma \text{(standard deviation)}$$

4.3.3 Learning Rate

$J(\theta)$ should decrease after every iteration. #iterations vary a lot.

Example *Automatic Convergence Test:* Declare convergence if $J(\theta)$ decreases by less than ϵ (say 10^{-3}) in one iteration.

If $J(\theta)$ increases, use smaller α. Too small α means slow convergence.

To choose α, try ..., 0.001, 0.003, 0.01, 0.03, 0.1, 0.3, 1, ...

4.4 Features and Polynomial Regression

4.4.1 Features

Get an insight in your problem and choose better features (may even combine/separate features).

Ex: size = length $\rightarrow$ breadth.

4.4.2 Polynomial Regression

Ex:

$$x_1 = size$$
$$x_2 = size^2$$
$$x_3 = size^3$$

4.5 Normal Equation

Solve for θ analytically!

$$x^{(i)} = \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix} \qquad \in \mathbb{R}^{n+1}$$

$$X = \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ \vdots \\ (x^{(m)})^T \end{bmatrix} \qquad \in \mathbb{R}^{m \times (n+1)}$$

$$= \begin{bmatrix} x_0 & x_1 & \dots & x_n \end{bmatrix} \qquad \in \mathbb{R}^{m \times (n+1)}$$

$$y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix} \qquad \in \mathbb{R}^m$$

$$\theta = (X^T X)^{-1} X^T y$$

Inverse of a matrix grows as $O(n^3)$, use wisely.

4.5.1 Non Invertibility of $X^T X$

Use 'pinv' function in Octave (pseudo-inverse) instead of 'inv' function (inverse).

If $X^T X$ is non-invertible, common causes are

1. Redundant features (linearly dependent)

2. Too many features ($m \leq n$). In this case, delete some features or use *regularization*

Check Lecture4.pdf for more details.

Chapter 5

Octave Tutorial

Check Lecture5.pdf for more details.

Chapter 6

Classification

Classify into categories (binary or multiple).

6.1 Logistic Regression

$$0 \leq h_\theta(x) \leq 1$$
$$h_\theta(x) = g(\theta^T x)$$
$$g(z) = \frac{1}{1 + e^{-z}} \qquad g \text{ is called a } \textit{sigmoid} \text{ function or a } \textit{logistic} \text{ function.}$$
$$h_\theta(x) = \frac{1}{1 + e^{\theta^T x}}$$

Interpretation of Hypothesis Output

$h_\theta(x) = $ estimated probability that $y = 1$ on input x

$h_\theta(x) = P(y = 1 | x; \theta) = $ probability that $y = 1$, given x, parameterized by θ

$$P(y = 0 | x; \theta) + P(y = 1 | x; \theta) = 1$$

6.2 Decision Boundary

$$\text{Predict: } y = 1 \qquad\qquad \text{if } h_\theta(x) \geq 0.5$$
$$(\theta^T x \geq 0)$$
$$\text{Predict: } y = 0 \qquad\qquad \text{if } h_\theta(x) < 0.5$$
$$(\theta^T x < 0)$$
$$\theta^T x = 0 \qquad\qquad \text{is the } \textit{decision boundary.}$$

Non-linear Decision Boundaries

Use same technique as polynomial regression for features.

6.3 Cost Function

$$\text{Training Set}: \{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), ./dots, (x^{(m)}, y^{(m)})\}$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$x_0 = 1$$
$$y \in \{0, 1\}$$
$$h_\theta(x) = \frac{1}{1 + e^{-\theta^T x}}$$

How to choose parameter θ?

Linear Regression:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^{m} \frac{1}{2} (h_\theta(x^{(i)}) - y^{(i)})^2$$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^{m} \text{Cost}(h_\theta(x^{(i)}), y^{(i)})$$

$$\text{Cost}(h_\theta(x), y) = \frac{1}{2} (h_\theta(x) - y)^2$$

Logistic Regression:

$$\text{Cost}(h_\theta(x), y) = \begin{cases} -\log(h_\theta(x)) & y = 1 \\ -\log(1 - h_\theta(x)) & y = 0 \end{cases}$$

$$\text{Cost}(h_\theta(x), y) = 0 \text{ if } h_\theta(x) = y$$
$$\text{Cost}(h_\theta(x), y) \to \inf \text{ if } y = 0 \text{ and } h_\theta(x) \to 1$$
$$\text{Cost}(h_\theta(x), y) \to \inf \text{ if } y = 1 \text{ and } h_\theta(x) \to 0$$

Note: $y = 0$ or 1 always.

$$\text{Cost}(h_\theta(x), y) = -y\log(h_\theta(x)) - (1 - y)\log(1 - h_\theta(x))$$

$$J(\theta) = \frac{1}{m}\sum_{i=1}^{m}\text{Cost}\left(h_\theta(x^{(i)}), y^{(i)}\right)$$

$$= -\frac{1}{m}\left[\sum_{i=1}^{m}\left(y^{(i)}\log(h_\theta(x^{(i)})) + (1 - y^{(i)})\log(1 - h_\theta(x^{(i)}))\right)\right]$$

To fit parameters θ:

$$\underset{\theta}{\text{minimize}} \; J(\theta)$$

To make a prediction given a new x:

$$\text{Output } h_\theta(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Gradient Descent:

Simultaneously update all θ_j

$$\theta_j := \theta_j - \alpha\frac{\partial}{\partial\theta_j}J(\theta)$$

Plug in the derivative

$$\theta_j := \theta_j - \alpha\frac{1}{m}\sum_{i=1}^{m}\left(h_\theta(x^{(i)}) - y^{(i)}\right)x_j^{(i)}$$

Don't forget feature scaling!

6.4   Advanced Optimization:

Something better than gradient descent:

1. Conjugate Gradient

2. BFGS

3. L-BFGS

Advantages:

1. No need to manually pick α

2. Often faster than gradient descent

Disadvantages:

1. More complex

Use libraries! Beware of bad implementations!

How to use: We first need to provide a function that evaluates the following two functions for a given input value of θ.

1. $J(\theta)$

2. $\dfrac{\partial}{\partial \theta_j} J(\theta)$

```
% We can write a single function that can return both of these:
function [jVal, gradient] = costFunction(theta)
        jVal = [...code to compute J(theta)...];
        gradient = [...code to compute derivative of J(theta)...];
end
```

Then we can use octave's *fminunc()* optimization algorithm along with the *optimset()* function that creates an object containing the options we want to send to *fminunc()*.

```
options = optimset('GradObj', 'on', 'MaxIter', 100);
initialTheta = zeros(2,1);
        [optTheta, functionVal, exitFlag] = fminunc(@costFunction,
                initialTheta, options);
```

6.5   Multi-Class Classification

6.5.1   One vs All

Build a separate binary classifier $h_\theta^{(i)}(x)$ for each class against all other classes.

$$h\theta^{(i)} = P(y = i | x; \theta) \ \forall i$$

On a new input x, to make a prediction, pick the class i that maximizes $h_\theta^{(i)}(x)$

Check Lecture6.pdf for more details.

Chapter 7

Regularization

7.1   Problem of Overfitting

1. Underfitting (High Bias)

2. Right Fit

3. Overfitting (High Variance)

Overfitting: If we have too many features, the learned hypothesis may fit the training set very well $\left(J(\theta) = \dfrac{1}{2m} \sum_{i=1}^{m} (h\theta(x^{(i)}) - y^{(i)})^2 \approx 0 \right)$, but fail to generalize to new examples (predict prices on new examples).

7.2   Addressing Overfitting

Options:

1. Reduce the number of features

   - Manually select which features to keep
   - Model selection algorithm (later)

2. Regularization

   - Keep all the features, but reduce the magnitude/values of parameters θ_j
   - Works well when we have a lot of features, each of which contributes a bit to predicting y

7.3 Cost Function

Say our overfitting hypothesis is $h_\theta(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$. Suppose, we penalize and make θ_3, θ_4 very small

$$\operatorname*{minimize}_{\theta} \left(\frac{1}{2m} \sum_{i=1}^{m} \left(h_\theta(x^{(i)}) - y^{(i)} \right)^2 + 1000\theta_3^2 + 1000\theta_4^2 \right)$$

Regularization

Small values for parameters $\theta_0, \theta_1, \ldots, \theta_n$.

- *Simpler* hypothesis

- Less prone to overfitting

Which parameters to penalize?

- Features: $x_1, x_2, \ldots, x_{100}$

- Parameters: $\theta_0, \theta_1, \theta_2, \ldots, \theta_{100}$

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^{m} \left(h_\theta(x^{(i)}) - y^{(i)} \right)^2 + \lambda \sum_{j=1}^{n} \theta_j^2 \right]$$

Note: The convention, we don't regularize θ_0 but it doesn't make very much difference.

λ here is *regularization parameter.*

What if λ is set too high (say 10^{10})? Underfitting!

7.4 Regularized Linear Regression

Updated $J(\theta)$:

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^{m} \left(h_\theta(x^{(i)}) - y^{(i)} \right)^2 + \lambda \sum_{j=1}^{n} \theta_j^2 \right]$$

$$\operatorname*{minimize}_{\theta} J(\theta)$$

Gradient Descent:

Repeat{

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^{m} (h_\theta(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \left[\frac{1}{m} \sum_{i=1}^{m} \left((h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{m} \theta_j \right] \quad \forall j \in \{1, 2, \ldots, n\}$$

$$\equiv \theta_j := \theta_j (1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^{m} (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}

Normal Equation:

$$X = \begin{bmatrix} \left(x^{(1)}\right)^T \\ \left(x^{(2)}\right)^T \\ \vdots \\ \left(x^{(m)}\right)^T \end{bmatrix} \qquad \in \mathbb{R}^{m \times (n+1)}$$

$$y = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix} \qquad \in \mathbb{R}^{m}$$

$$\theta = \left(X^T X + \lambda \begin{bmatrix} 0 & 0 & 0 & \ldots & 0 \\ 0 & 1 & 0 & \ldots & 0 \\ 0 & 0 & 1 & \ldots & 0 \\ 0 & 0 & 0 & \ddots & 0 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix} \left(\in \mathbb{R}^{(n+1) \times (n+1)} \right) \right)^{-1} - X^T y \quad \in \mathbb{R}^{n+1}$$

Non-Invertibility

Suppose $m \leq n$,

$$\theta = \left(X^T X \right)^{-1} X^y$$

If $\lambda > 0$,

$$\theta = \left(X^T X + \lambda \begin{bmatrix} 0 & & & & \\ & 1 & & & \\ & & 1 & & \\ & & & \ddots & \\ & & & & 1 \end{bmatrix} \right)^{-1} - X^T y$$

Not a problem!

7.5   Regularized Logistic Regression

Cost Function:

$$J(\theta) = - \left[\frac{1}{m} \sum_{i=1}^m \left(y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_\theta(x^{(i)})) \right) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Gradient Descent:

Repeat{

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m \left((h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{m} \theta_j \right] \qquad \forall j \in \{1, 2, \dots, n\}$$

$$\equiv \theta_j := \theta_j (1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}

Advanced Optimization:

```
function [jVal, gradient] = costFunction(theta)
        jVal = [...code to compute J(theta)...];
        gradient = [...code to compute derivative of J(theta)...];
```

Check Lecture7.pdf for more details.

Chapter 8

Neural Networks

8.1   Non-Linear Hypothesis

If #features is high, hypothesis could have extremely high #terms. Hence the need for non-linear hypothesis.

8.2   Neural Networks

- *Origins:* Algorithms that try to mimic brain.

- Widely used in 80s and early 90s; diminished in late 90s.

- Resurgance: State-of-the-art technique for many applications.

8.3   Model Representation

Neuron Model: Logistic Unit

1. input layer

2. hidden layer / computation layer

3. output layer

4. activation function ($g()$)

parameters $\equiv$ weights

8.4 Notations

$$a_i^{(j)} = activation \text{ of unit } i \text{ in layer } j$$
$$\Theta^{(j)} = \text{matrix of weights controlling}$$
$$\text{function mapping from layer } j \text{ to}$$
$$\text{layer } j+1$$

Example:

$$\text{layer } 1 : \{x_1, x_2, x_3\}$$
$$\text{layer } 2 : \{a_1^{(2)}, a_2^{(2)}, a_3^{(2)}\}$$
$$\text{layer } 3 : \{a_1^{(3)}\}$$

$$a_1^{(2)} = g\left(\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3\right)$$
$$a_2^{(2)} = g\left(\Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3\right)$$
$$a_3^{(2)} = g\left(\Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3\right)$$
$$h_\Theta(x) = a_1^{(3)} = g\left(\Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(1)} a_3^{(2)}\right)$$

Note: If network has s_j units in layer j, and s_{j++1} units in layer $j+1$, then

$$\Theta^{(j)} \in \mathbb{R}^{s_{j+1} \times (s_j+1)}$$

8.5 Forward Propagation

$$z_1^{(2)} = \Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3$$

$$z_2^{(2)} = \Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3$$

$$z_3^{(2)} = \Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3$$

$$a_1^{(2)} = g(z_1^{(2)})$$

$$a_2^{(2)} = g(z_2^{(2)})$$

$$a_3^{(2)} = g(z_3^{(2)})$$

Vectorized Implementation

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 3 \end{bmatrix}, \ x_0 = 1$$

$$z^{(2)} = \begin{bmatrix} z_1^{(2)} \\ z_2^{(2)} \\ z_3^{(2)} \end{bmatrix}$$

$$a^{(1)} = x \qquad\qquad \in \mathbb{R}^4$$

$$z^{(2)} = \Theta^{(1)} a^{(1)} \qquad\qquad \in \mathbb{R}^3$$

$$a^{(2)} = g(z^{(2)}) \qquad\qquad \in \mathbb{R}^3$$

$$\textbf{Add } a_0^{(2)} = 1 \qquad\qquad \to a^{(2)} \in \mathbb{R}^4$$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

$$h_\Theta(x) = a^{(3)} = g(z^{(3)})$$

Other network architechtures are possible!

Non-linear classification example: XOR/XNOR

x_1, x_2 are binary (0 or 1)

$$y = x_1 \text{ XOR } x_2 \qquad\qquad\qquad or$$
$$= x_1 \text{ XNOR } x_2 \qquad\qquad \equiv \text{NOT}(x_1 \text{ XOR } x_2)$$

[1] Simple example: AND

$$x_0 = 0$$
$$x_1, x_2 \in \{0, 1\}$$
$$y = x_1 \text{ AND } x_2$$
$$\Theta_{10}^{(1)} = -30$$
$$\Theta_{11}^{(1)} = 20$$
$$\Theta_{12}^{(1)} = 20$$
$$h_\Theta(x) = g(-30 + 20x_1 + 20x_2)$$

x_1	x_2	$h_\Theta(x)$
0	0	$g(-30) \approx 0$
0	1	$g(-10) \approx 0$
1	0	$g(-10) \approx 0$
1	1	$g(10) \approx 1$

We can combine AND, OR, NOT, (NOT x_1) AND (NOT x_2) to get XNOR, XOR, etc.

8.6   Multi-Class Classification

Extension of One-vs-All Method!

Example: Say 4 classes, so 4 output units. So, $y \in \mathbb{R}^4$ with 1 in corresponding class and 0 elsewhere.

---

[1]Work out more examples (OR, NOT, (NOT x_1) AND (NOT x_2)!
  Check Lecture8.pdf for more details.

Chapter 9

Back Propagation

9.1 Notations

$$L = \text{total no. of layers in the network}$$
$$s_l = \text{no. of units (not counting bias unit) in layer } l$$

Binary Classification

$$y = 0 \text{ or } 1$$

1 output unit

$$h_\Theta(x) \in \mathbb{R}$$
$$s_L = 1$$
$$K = 1 \qquad \text{(for simplification)}$$

Multi-class Classification (K classes)

$$y \in \mathbb{R}^K$$

Example: $K = 4$

$$y \in \left\{ \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix} \right\}$$

K output units

$$h_\Theta(x) \in \mathbb{R}^K$$
$$s_L = K$$
$$K \geq 3$$

9.2 Cost Function

Logistic Regression

$$J(\theta) = -\left[\frac{1}{m}\sum_{i=1}^{m}\left(y^{(i)}\log h_\theta(x^{(i)}) + (1-y^{(i)})\log(1-h_\theta(x^{(i)}))\right)\right] + \frac{\lambda}{2m}\sum_{j=1}^{n}\theta_j^2$$

Neural Network

$$h_\Theta(x) \in \mathbb{R}^K \quad (h_\Theta(x))_i = i^{th}\text{output}$$

$$J(\Theta) = -\frac{1}{m}\left[\sum_{i=1}^{m}\sum_{k=1}^{K}\left(y_k^{(i)}\log\left(h_\Theta(x^{(i)})\right)_k + \left(1-y_k^{(i)}\right)\log\left(1-\left(h_\Theta(x^{(i)})\right)_k\right)\right)\right]$$
$$+ \frac{\lambda}{2m}\sum_{l=1}^{L-1}\sum_{i=1}^{s_l}\sum_{j=1}^{s_{l+1}}\left(\Theta_{ji}^{(l)}\right)^2$$

9.3 Backpropagation Algorithm

9.3.1 Gradient Computation

$$J(\Theta) = -\frac{1}{m}\left[\sum_{i=1}^{m}\sum_{k=1}^{K}\left(y_k^{(i)}\log\left(h_\Theta(x^{(i)})\right)_k + \left(1-y_k^{(i)}\right)\log\left(1-\left(h_\Theta(x^{(i)})\right)_k\right)\right)\right]$$
$$+ \frac{\lambda}{2m}\sum_{l=1}^{L-1}\sum_{i=1}^{s_l}\sum_{j=1}^{s_{l+1}}\left(\Theta_{ji}^{(l)}\right)^2$$

$$\min_\Theta J(\Theta)$$

Need code to compute:

- $J(\Theta)$

- $\dfrac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$

Given one training example (x, y):

Forward Propagation:

$$
\begin{aligned}
a^{(1)} &= x \\
z^{(2)} &= \Theta^{(1)} a^{(1)} \\
a^{(2)} &= g(z^{(2)}) \qquad\qquad &\text{(add } a_0^{(2)}\text{)} \\
z^{(3)} &= \Theta^{(2)} a^{(2)} \\
a^{(3)} &= g(z^{(3)}) \qquad\qquad &\text{(add } a_0^{(3)}\text{)} \\
z^{(4)} &= \Theta^{(3)} a^{(3)} \\
h_\Theta(x) = a^{(4)} &= g(z^{(4)})
\end{aligned}
$$

Gradient Computation: Back Propagation

Intuition: $\delta_j^{(l)} =$ "error" of node j in layer l

For each output unit ($L = 4$)

$$
\delta_j^{(4)} = a_j^{(4)} - y_j
$$

Vectorize:

$$
\begin{aligned}
\delta^{(4)} &= a^{(4)} - y \\
\delta^{(3)} &= (\Theta^{(3)})^T \delta^{(4)} .* g'(z^{(3)}) \\
g'(z^{(3)}) &= a^{(3)} .* (1 - a^{(3)}) \\
\delta^{(2)} &= (\Theta^{(2)})^T \delta^{(3)} .* g'(z^{(2)}) \\
g'(z^{(2)}) &= a^{(2)} .* (1 - a^{(2)})
\end{aligned}
$$

No $\delta^{(1)}$

$$
\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = a_j^{(l)} \delta_i^{(l+1)} \quad \text{(ignoring } \lambda; \text{ if } \lambda = 0)
$$

Backpropagation Algorithm

Training set $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$

Set $\Delta_{ij}^{(l)} = 0$ (for all l, i, j) $\rightarrow$ accumulators to compute $\dfrac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$

For $i = 1$ to m

   Set $a^{(1)} = x^{(1)}$

   Perform forward propagation to compute $a^{(l)}$ for $l = 2, 3, \dots, L$

   Using $y^{(i)}$, compute $\delta^{(L)} = a^{(L)} - y^{(i)}$

   Compute $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$

   $\Delta_{ij}^{(l)} := \Delta_{ij}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$

   $\Delta^{(l)} := \Delta^{(l)} + \delta^{(l+1)} (a^{(l)})^T$ (Vectorized)

$$D_{ij}^{(l)} := \begin{cases} \dfrac{1}{m} \Delta_{ij}^{(l)} + \lambda \Theta_{ij}^{(l)} & \text{if } j \neq 0 \\[2em] \dfrac{1}{m} \Delta_{ij}^{(l)} & \text{if } j = 0 \end{cases}$$

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = D_{ij}^{(l)}$$

9.4 Implementation Note

9.4.1 Unrolling Parameters

Advanced optimization functions like 'fminunc' assume that inputs 'theta', 'gradient', etc. are vectors. These are matrices in Neural Networks, so we *unroll* them into vectors.

Example: Neural Network with 3 layers (1 input, 1 hidden, 1 output) and $s_1 = 10$, $s_2 = 10$, $s_3 = 1$.

$$\Theta^{(1)} \in \mathbb{R}^{10 \times 11}, \ \Theta^{(2)} \in \mathbb{R}^{10 \times 11}, \ \Theta^{(3)} \in \mathbb{R}^{(1 \times 11)}$$
$$D^{(1)} \in \mathbb{R}^{10 \times 11}, \ D^{(2)} \in \mathbb{R}^{10 \times 11}, \ D^{(3)} \in \mathbb{R}^{(1 \times 11)}$$

To unroll:

```
thetaVec = [Theta1(:); Theta2(:); Theta3(:)];
DVec = [D1(:); D2(:); D3(:)];
```

To go back to matrices:

```
Theta1 = reshape(thetaVec(1:110), 10, 11);
Theta2 = reshape(thetaVec(111:220), 10, 11);
Theta3 = reshape(thetaVec(221:231), 1, 11);
```

9.4.2   Learning Algorithm

Have initial parameters $\Theta^{(1)}, \Theta^{(2)}, \Theta^{(3)}$

Unroll to get `initialTheta` to pass to

`fminunc(@costFunction, initialTheta, options)`

`function [jVal, gradientVec] = costFunction(thetaVec)`

From `thetaVec`, get $\Theta^{(1)}, \Theta^{(2)}, \Theta^{(3)}$ using `reshape`

Use forward/back propagation to compute $D^{(1)}, D^{(2)}, D^{(3)}$ and $J(\Theta)$

Unroll $D^{(1)}, D^{(2)}, D^{(3)}$ to get `gradientVec`

9.5   Gradient Checking

An unfortunate problem is that it might seem like it's working even when there's a bug in the backpropagation algorithm (implementation). Subtle bugs can hence cause higher level of error and perform worse than a bug-free implementation. Here, we try to make sure that our implementation is 100% correct.

9.5.1   Numerical Estimation of gradients

To estimate derivative of $J(\theta)$, we calculate $J(\theta + \epsilon)$ and $J(\theta - \epsilon)$ and take the slope of the line connecting these two points.

$$J'(\theta) \approx \frac{J(\theta + \epsilon) - J(\theta - \epsilon)}{2\epsilon}$$

Use pretty small ϵ ($\approx 10^{-4}$).

The above formula is called the two sided difference. Another formula called the one sided difference is

$$J'(\theta) \approx \frac{J(\theta + \epsilon) - J(\theta)}{\epsilon}$$

But the two sided difference usually gives a slightly more accurate approximation.

Implement:

```
gradApprox = (J(theta + EPSILON) - J(theta - EPSILON)) / (2*EPSILON)
```

Parameter *vector* θ

$$\theta \in \mathbb{R}^n$$

$$\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix}$$

$$\frac{\partial}{\partial \theta_1} J(\theta) \approx \frac{J(\theta_1 + \epsilon, \theta_2, \theta_3, \ldots, \theta_n) - J(\theta_1 - \epsilon, \theta_2, \theta_3, \ldots, \theta_n)}{2\epsilon}$$

$$\frac{\partial}{\partial \theta_2} J(\theta) \approx \frac{J(\theta_1, \theta_2 + \epsilon, \theta_3, \ldots, \theta_n) - J(\theta_1, \theta_2 - \epsilon, \theta_3, \ldots, \theta_n)}{2\epsilon}$$

$$\vdots$$

$$\frac{\partial}{\partial \theta_n} J(\theta) \approx \frac{J(\theta_1, \theta_2, \theta_3, \ldots, \theta_n + \epsilon) - J(\theta_1, \theta_2, \theta_3, \ldots, \theta_n - \epsilon)}{2\epsilon}$$

Implement:

```
for i = 1:n
        thetaPlus = theta;
        thetaPlus(i) = thetaPlus(i) + EPSILON;
        thetaMinus = theta;
        thetaMinus(i) = thetaMinus(i) - EPSILON;
        gradApprox(i) = (J(thetaPlus) - J(thetaMinus))/(2*EPSILON);
end;
```

Check that `gradApprox` $\approx$ `DVec` (from backpropagation).

9.5.2   Implementation Note:

- Implement backpropagation to compute `Dvec` (unrolled $D^{(1)}, D^{(2)}, \ldots$)

- Implement numerical gradient check to compute `gradApprox`

- Make sure they give similar values

- Turn off gradient checking (much slower) and use backpropagation for learning (much faster)

Important:

- Be sure to disable your gradient checking code before training your classifier. If you run numerical gradient computaion on every iteration of gradient descent (or in the inner loop of `costFunction(...)`) your code will be <u>very</u> slow.

9.6   Random Initialization

Initial value of Θ

For gradient descent and advanced optimization methods, we need initial value of Θ.

Consider gradient descent:

Set `initialTheta = zeroes(n, 1)`?

Works fine for linear/logistic regression, but not when we're training a neural network. Why?

$$\Theta_{ij}^{(l)} = 0 \text{ for all } i, j, l$$
$$a_1^{(2)} = a_2^{(2)}$$
$$\delta_1^{(2)} = \delta_2^{(2)}$$
$$\frac{\partial}{\partial \Theta_{01}^{(1)}} J(\Theta) = \frac{\partial}{\partial \Theta_{02}^{(1)}} J(\Theta)$$

Meaning, after each gradient descent update, parameters corresponding to inputs going into each unit of the next layer are identical. Thus, all the units of the next layer are still computing the same function of the input. The values may change, but will always be the same, preventing the neural network from learning something interesting since it is extremely redundant.

Random Initialization: Symmetry Breaking

To get around this (*symmetric weights*), we do random initialization. Initialize each $\Theta_{ij}^{(l)}$ to a random value in $[-\epsilon, \epsilon]$ (i.e. $-epsilon \leq \Theta_{ij}^{(l)} \leq \epsilon$).

Example:

```
Theta1 = rand(10,11)*(2*INIT_EPSILON) - INIT_EPSILON
Theta2 = rand(1, 11)*(2*INIT_EPSILON) - INIT_EPSILON
```

A good choice of ϵ is:

$$\epsilon = \frac{\sqrt{6}}{\sqrt{L_{in} + L_{out}}}$$

Where L_{in} and L_{out} are number of units in the layers adjacent to $\Theta^{(l)}$.

9.7   Putting it all together

Pick a network architecture (connectivity pattern between neurons), like number of layers, number of units in each layer, etc. For example, $[3 \rightarrow 5 \rightarrow 4]$ vs $[3 \rightarrow 5 \rightarrow 5 \rightarrow 4]$ vs $[3 \rightarrow 5 \rightarrow 5 \rightarrow 5 \rightarrow 4]$.

$$\text{No. of input units} : \text{Dimension of features } x^{(i)}$$
$$\text{No. of output units} : \text{Number of classes}$$

Remember:

$$y \in \{1, 2, \ldots, 10\} \text{ becomes}$$

$$y \in \left\{ \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}, \ldots, \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix} \right\}$$

For Hidden Layers?

Reasonable Default: 1 hidden layer, or if > 1 hidden layer, have same no. of hidden units in every layer (usually, the more the better, just that it may get computationally expensive). Usually, comparable to the number of features.

Training a neural network:

1. Randomly initialize weights

2. Implement forward propagation to get $h_\Theta(x)$ for any x

3. Implement code to compute cost function $J(\Theta)$

4. Implement backpropagation to compute partial derivatives $\dfrac{\partial}{\partial\Theta_{ij}^{(l)}}J(\Theta)$

```
for i = 1:m
        Perform forward propagation ...
                and backpropagation ...
                using example (x⁽ⁱ⁾, y⁽ⁱ⁾)

        Get activations a⁽ˡ⁾ ...
                and delta terms δ⁽ˡ⁾ ...
                for l = 2,...,L

        Δ⁽ˡ⁾ := Δ⁽ˡ⁾ + δ⁽ˡ⁺¹⁾(a⁽ˡ⁾)ᵀ
        ⋮
end
⋮
compute  ∂/∂Θ⁽ˡ⁾_jk J(Θ)
```

5. (a) Use gradient checking to compare $J'(\Theta)$ computed using backpropagation vs using numerical estimate.

   (b) Then disable gradient checking code.

6. Use gradient descent or advanced optimization method with backpropagation to try to minimize $J(\Theta)$

$J(\Theta)$ in general for a neural network is non-convex, and hence susceptible to local optima. However, it turns out, in practice, this is not usually a huge problem.

Check Lecture9.pdf for more details.

Chapter 10

Applying Machine Learning

10.1 Deciding what to try next

10.1.1 Debugging a Learning Algorithm

Suppose you have implemented regularized linear regression to predict housing prices.

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^{m} \left(h_\theta(x^{(i)}) - y^{(i)} \right)^2 + \lambda \sum_{j=1}^{n} \theta_j^2 \right]$$

However, it makes *unacceptably* large errors in prediction. What to try next?

- Get more training examples (Sometimes doesn't actually help)

- Try smaller set of features (Overfitting?)

- Try additional features (Underfitting?) (Huge project)

- Try adding polynomial features

- Increasing/Decreasing λ

Need to understand what would help.

10.1.2 Machine Learning Diagnostics

Diagnostic: A test that you can run to gain insight at what is/isn't working with the algorithm, and gain guidance as to how best to improve its performance.

Diagnostics can take time to implement, but doing so can be a very good use of your time as it can save huge time later.

10.2   Evaluating your Hypothesis

Just *training error* not a good indicator (Overfitting).

Solution? Plot Hypothesis function? Not feasible with high number of features.

10.2.1   The Standard Way

Split the dataset into two portions - *Training Set* and *Test Set* with about 70%-30% ratio (*randomly* if it is in some order).

Notation:

$$\text{Training Examples}: (x, y)$$
$$\text{Test Examples}: (x_{test}, y_{test})$$
$$\text{No. of Training Examples}: m$$
$$\text{No. of Test Examples}: m_{test}$$

10.2.2   Train/Test Procedure

1. Learn θ from training data (minimize $J(\theta)$)

2. Compute test set error ($J_{test}(\theta)$) on test examples

   - In case of classification problem, we can also use misclassification error (Also called 0/1 misclassification error)

   - $\text{err}(h_\theta(x), y) = \begin{cases} 1 & \text{if } h_\theta(x) \geq 0.5, \ y == 0 \\ & \text{or if } h_\theta(x) < 0.5, \ y == 1 \\ 0 & \text{otherwise} \end{cases}$

   - $\text{test error} = \dfrac{1}{m_{test}} \sum_{i=1}^{m_{test}} \text{err}(h_\theta(x_{test}^{(i)}), y_{test}^{(i)})$

   - This gives us the proportion of the test data that was misclassified

10.3   Model Selection

Once parameters θ were fit to some training set, the error $J(\theta)$ measured on that data is likely to be lower than the actual generalization error.

Model Selection: Let's say we try to choose what degree polynomial to fit to data (linear, quadradic, cubic, …)

Notation:

$$d = \text{Degree of polynomial}$$

For different d, we can have the following:

1. $h_\theta(x) = \theta_0 + \theta_1 x$

2. $h_\theta(x) = \theta_0 + \theta_1 x + \theta_2 x^2$

3. $h_\theta(x) = \theta_0 + \theta_1 x + \cdots + \theta_3 x^3$
    $\vdots$

10. $h_\theta(x) = \theta_0 + \theta_1 x + \cdots + \theta_{10} x^{10}$

Choose one model out of the above. We can train them separately, and get different $\theta^{(k)}$, and calculate $J_{test}(\theta)$ for each of them, and choose the one with the least test error.

How well does this model generalize? Report $J_{test}(\theta)$.

Problem: $J_{test}(\theta)$ is likely an optimistic estimate of generalization error (our extra parameter d is fit to *test set*).

10.3.1   Evaluating your hypothesis

Divide the data set in three pieces, *training set* (60%), *cross-validation set* (20%), *test set* (20%).

Notations:

$$\text{Training Examples} : (x, y)$$
$$\text{Cross-Validation Examples} : (x_{cv}, y_{cv})$$
$$\text{Test Examples} : (x_{test}, y_{test})$$
$$\text{No. of Training Examples} : m$$
$$\text{No. of Cross-Validation Examples} : m_{cv}$$
$$\text{No. of Test Examples} : m_{test}$$
$$\text{Training Error} : J(\theta) \text{ or } J_{train}(\theta)$$
$$\text{Cross Validation Error} : J_{cv}(\theta)$$
$$\text{Test Error} : J_{test}(\theta)$$

Use *Cross-Validation Set* to select the model. Estimate generalization error for test set.

10.4 Bias vs Variance

Plot *error* against *degree* of the polynomial for both, *training* and *validation* data set. Usually, training error tends to decrease. Cross-Validation error (or test error) on the other hand, tends to decrease, and then increase again.

To recognise bias/variance problem:

$$\text{Bias (Underfit)} : \qquad J_{train} \text{ will be high}$$
$$J_{cv} \approx J_{train}$$
$$\text{Variance (Overfit)} : \qquad J_{train} \text{ will be low}$$
$$J_{cv} >> J_{train}$$

10.4.1 Effect of regularization

Large λ can cause high bias (underfit), and at very small λ causes high variance (overfit). How to automatically choose a good value of λ.

Choosing λ:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^{m} (h_\theta(x^{(i)}) - y^{(i)}) + \frac{\lambda}{2m} \sum_{j=1}^{m} \theta_j^2$$

$$J_{train}(\theta) = \frac{1}{2m} \sum_{i=1}^{m} (h_\theta(x^{(i)}) - y^{(i)})$$

$$J_{cv}(\theta) = \frac{1}{2m_{cv}} \sum_{i=1}^{m_{cv}} (h_\theta(x_{cv}^{(i)}) - y_{cv}^{(i)})$$

$$J_{test}(\theta) = \frac{1}{2m_{test}} \sum_{i=1}^{m_{test}} (h_\theta(x_{test}^{(i)}) - y_{test}^{(i)})$$

Have a model ($h_\theta(x)$, $J(\theta)$), have a bunch of values to try for λ like 0, 0.01, 0.02, 0.04, 0.08, ..., 10 (roughly doubling). Train (minimize $J(\theta)$) for all values. Use $J_{cv}(\theta)$ to choose a λ. Report generalization error as $J_{test}(\theta)$.

Bias/Variance as a function of λ

Plot $J_{train}(\theta)$ and $J_{cv}(\theta)$ as a function of λ. Usually, $J_{train}(\theta)$ tends to increase while $J_{cv}(\theta)$ tends to decrease at first but increase again later.

10.5 Learning Curves

Useful thing to plot, to sanity check, to improve the performance, to diagnose for bias/variance, ...

Plot $J_{train}(\theta)$, and $J_{cv}(\theta)$ as a function of m, the number of training examples. Measure $J_{train}(\theta)$ only on data used to train!. Measure $J_{cv}(\theta)$ on all cross validation data!

Usually, average training error will increase with m. Usually, average validation error will decrease with m.

10.5.1 High Bias

Usually: Validation error will decrease at first but then will saturate/flatten out soon. Training error will increase at first but then will saturate/flatten out soon, close to validation error.

Note: If a learning algorithm is suffering from high bias, getting more training data will not (by itself) help much.

10.5.2 High Variance

Usually: Training error will increase but still be pretty low. Validation error will decrease but still be pretty high, far from training error.

Note: If a learning algorithm is suffering from high variance, getting more training examples is indeed likely to help.

10.6 Deciding what to try next (Revisited)

Get more training examples	fixes high variance
Try smaller set of features	fixes high variance
Try additional features	fixes high bias (not always)
Try adding polynomial features	fixes high bias (not always)
Increasing/Decreasing λ	fixes bias/variance

10.6.1 Neural Networks

Small neural network, computationally cheaper, but possibly underfitting. *Large* neural networks, computationally more expensive, and possibly overfitting. Use regularization to address overfitting.

Often, large neural network with regularization works better (but may be computationally expensive).

Large network can comprise of a single large hidden layer, or two moderate hidden layers, or three average hidden layers, etc.

Check Lecture10.pdf for more details.

Chapter 11

Machine Learning System Design

11.1 Prioritizing what to work on: Spam Classification Example

Supervised Learning! x = features of email. y = spam(1) or not spam(0).
Features x: Choose 100 words indicative of spam/not spam.
Example: deal, buy, discount, andrew, now, ...

Feature vector:

$$x = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \\ \vdots \\ 1 \end{bmatrix} \begin{bmatrix} andrew \\ buy \\ deal \\ discount \\ \vdots \\ now \end{bmatrix}$$

0/1 depending on whether a 'word' appears or not in the email.

$$x_j = \begin{cases} 1 & \text{if word } j \text{ appearsin email} \\ 0 & \text{otherwise} \end{cases}$$

Note: In practice, take most frequently occurring $n(10,000 \text{ to } 50,000)$ words in training set, instead of manually picking 100 words.

11.1.1 Time Management?

- Collect lots of data

– E.g. *honeypot* project.

- Develop sophisticated features based on email routing information (from email header).

- Develop sophisticated features for message body, e.g. should *discount* and *discounts* be treated as same word? How about *deal* and *Dealer*? Features about punctuation?

- Develop sophisticated algorithm to detect misspellings (e.g. m0rtgage, med1cine, w4tches.)

11.2 Error Analysis

11.2.1 Recommended approach

- Start with a simple and quick algorithm, implement and test it on cross-validation data.

- Plot learning curves to decide if more data, more features, etc.

- Error Analysis: Manually examine the examples (in cross validation set) that your algorithm made errors on. See if you spot any systematic trend in what type of examples it is making errors on.

m_{cv} = 500 examples in cross validation set, algorithm misclassifies 100 emails. Manually examine 100 errors, and categorized them based on:

1. What type of email it is (pharma, replica, phishing, ...)

2. What features would have helped (Deliberate misspellings, Unusual email routing, Unusual punctuation)

11.2.2 Numerical Evaluation

Have a way of *numerical evaluation* like accuracy/error, to tell how the learning algorithm is doing.

Example: Should discount/discounts/discounting be treated as the same word? Can use *stemming* software (E.g. "Porter Stemmer"). It can hurt because the software can mistake universe/university to be the same.

So, error analysis may not be helpful for deciding if this is likely to improve performance. Only solution is to try and see if it works.

Need numerical evaluation (e.g. cross validation error) of algorithm's performance with and without stemming.

11.3 Skewed Classes

Cancer classification example: Train logistic regression model $h_\theta(x)$ with 1% error on test set. Good? What if only 0.50% patients in training/test set have cancer? Not so impressive anymore! Even predicting *No Cancer* all the time will get 0.50% error. This is a case of *skewed* classes.

Going from 99.2% to 99.5% accuracy is actually doing something useful? Can't say in case of skewed classes.

11.3.1 Precision/Recall

$y = 1$ in presence of rare class that we want to detect.

Notations

Notations	Actual Class	Predicted Class
True Positive	1	1
True Negative	0	0
False Positive	0	1
False Negative	1	0

Precision

Out of all the predicted *positives*, what fraction is *actually* positive.

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

Recall

Out of all the actual *positives*, what fraction is *predicted* positive.

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

Now, if we predict no cancer always, then recall $= 0$. These metrics usually help us evaluate our algorithm even in case of really skewed classes.

11.3.2 Trade-Off

Consider same example of cancer, we train logistic regression $h_\theta(x)$ which gives the probability of having cancer.

Suppose we want to predict cancer ($y = 1$) only if we're very sure. One way to do this, is predict $y = 1$ when $h_\theta(x) \geq 0.7$ instead of 0.5 and $y = 0$ otherwise. This will give us *higher* precision, but *lower* recall.

Suppose we want to avoid missing too many cases of cancer, we can predict $y = 1$ when $h_\theta(x) \geq 0.3$. This will give us *higher* recall, but *lower* precision.

How to choose the threshold?

F_1 Score (F Score)

We have lost our 'single number' evaluation metric, we now have precision and recall. What are our options?

- Average of precision and recall ($\dfrac{P + R}{2}$)? Not a good idea! (Extremes)

- F_1 Score $= 2\dfrac{PR}{P + R}$ better choice, commonly used

Use F_1 score on cross-validation set to decide the threshold.

11.3.3 Working with Large Data

Consider cases where more data will be helpful.

Say, feature $x \in \mathbb{R}^{n+1}$ has sufficient information to predict y accurately. (Counter-examples: predict housing price from only size or area).

A useful test is, given x, can a human expert confidently predict y?

Alongside, Suppose we use a learning algorithm with a large number of parameters (low bias algorithms).

Using a very large training set, it is unlikely to overfit.

Check Lecture11.pdf for more details.

Chapter 12

Support Vector Machines

Another popular supervised learning algorithm.

12.1 Optimization Objective

12.1.1 Alternative view of logistic regression

$$h_\theta(x) = \frac{1}{1 + e^{-\theta^T x}}$$

If $y = 1$, we want $h_\theta(x) \approx 1$, $\theta^T x \gg 0$
If $y = 0$, we want $h_\theta(x) \approx 0$, $\theta^T x \ll 0$

Cost of example:

Each example (x, y) contributes:

$$-(y \log(h_\theta(x)) + (1 - y) \log(1 - h_\theta(x)))$$
$$= -y \log\left(\frac{1}{1 + e^{\theta^T x}}\right) - (1 - y) \log\left(1 - \frac{1}{1 + e^{\theta^T x}}\right)$$

If $y = 1$ (want $\theta^T x \gg 0$): we get $-log\left(\frac{1}{1 + e^{-z}}\right)$ where $z = \theta^T x$. For *SVM* modify cost function to be flat 0 for $x > 1$ and a straight line to the left (with negative slope). Denote by $cost_1(z)$.

If $y = 0$ (want $\theta^T x \ll 0$): we get $-log\left(1 - \dfrac{1}{1 + e^{-z}}\right)$ where $z = \theta^T x$. For *SVM* modify cost function to be flat 0 for $x < -1$ and a straight line to the left (with positive slope). Denote by $cost_0(z)$.

Logistic Regression:

$$\min_{\theta} \frac{1}{m}\left[\sum_{i=1}^{m}\left(y^{(i)}\left(-\log\left(h_\theta(x^{(i)})\right)\right) + (1 - y^{(i)})\left(-log\left(1 - h_\theta(x^{(i)})\right)\right)\right)\right] + \frac{\lambda}{2m}\sum_{j=1}^{n}\theta_j^2$$

Support Vector Machine:

Modify costs first

$$\min_{\theta} \frac{1}{m}\left[\sum_{i=1}^{m}\left(y^{(i)}cost_1(\theta^T x^{(i)}) + (1 - y^{(i)})cost_0(\theta^T x^{(i)})\right)\right] + \frac{\lambda}{2m}\sum_{j=1}^{n}\theta_j^2$$

Get rid of $\dfrac{1}{m}$ term. Only a constant, same optimal value!

$$\min_{\theta}\left[\sum_{i=1}^{m}\left(y^{(i)}cost_1(\theta^T x^{(i)}) + (1 - y^{(i)})cost_0(\theta^T x^{(i)})\right)\right] + \frac{\lambda}{2}\sum_{j=1}^{n}\theta_j^2$$

By convention, instead of optimizing $A + \lambda B$, optimize $CA + B$ where C is somewhat equivalent of $\dfrac{1}{\lambda}$, again, same optimal value!

$$\min_{\theta} C\left[\sum_{i=1}^{m}\left(y^{(i)}cost_1(\theta^T x^{(i)}) + (1 - y^{(i)})cost_0(\theta^T x^{(i)})\right)\right] + \frac{1}{2}\sum_{j=1}^{n}\theta_j^2$$

12.1.2   SVM Hypothesis

$$\min_{\theta} C\left[\sum_{i=1}^{m}\left(y^{(i)}cost_1(\theta^T x^{(i)}) + (1 - y^{(i)})cost_0(\theta^T x^{(i)})\right)\right] + \frac{1}{2}\sum_{j=1}^{n}\theta_j^2$$

Hypothesis:

Unlike logistic regression, SVM does not give probability. It just makes prediction.

$$h_\theta(x) = \begin{cases} 1 & \text{if } \theta^T x \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

12.2   Large Margin

12.2.1   Intuition

If $y = 1$, we want $\theta^T x \geq 1$ (not just ≥ 0)
If $y = 0$, we want $\theta^T x \leq -1$ (not just < 0)
This builds safety margin factor in between.

SVM Decision Boundary

If $C \gg 1$, say $100,000$, we will get
$$\min\left(C \times 0 + \frac{1}{2}\sum_{j=1}^{n}\theta_j^2\right) \quad \text{such   that} \quad \begin{cases} \theta^T x^{(i)} \geq 1 & \text{if } y^{(i)} = 1 \\ \theta^T x^{(i)} \leq -1 & \text{if } y^{(i)} = 0 \end{cases} \quad \text{Very   interesting}$$
decision boundary!

Linearly Separable Case: It will choose the decision boundary such that it maximises margin of the SVM.

Outliers: If C is indeed very large it will be affected by outliers easily but if C is not too large, it will still give a reasonable boundary.

12.2.2   Math

Vector Inner Product

$u^T v$ (or dot product)

$$u = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix}$$

$$v = \begin{bmatrix} v_1 \\ v_2 \end{bmatrix}$$

$$\|u\| = \text{length of vector } u$$

$$= \sqrt{u_1^2 + u_2^2} \in \mathbb{R}$$

$$p = \text{projection of vector } v \text{ on vector } u$$

$$u^T v = p\,\|u\|$$

$$u^T v = u_1 v_1 + u_2 v_2$$

$$u^T v = v^T u$$

SVM Decision Boundary

$$\min_{\theta} \frac{1}{2} \sum_{j=1}^{n} \theta_j^2 \text{ s.t. } \begin{cases} \theta^T x^{(i)} \geq 1 & \text{if } y^{(i)} = 1 \\ \theta^T x^{(i)} \leq -1 & \text{if } y^{(i)} = 0 \end{cases}$$

Simplifications: Only for the purpose of understanding here!

$$\theta_0 = 0$$
$$n = 2$$

This gives, $\min_{\theta} \frac{1}{2}(\theta_1^2 + \theta_2^2) = \min_{\theta} \frac{1}{2}\left(\sqrt{\theta_1^2 + \theta_2^2}\right)^2 = \min_{\theta} \frac{1}{2} \|\theta\|^2$

$\theta^T x^{(i)} = ?$

$$p^{(i)} = \text{projection of vector } x^{(i)} \text{ on vector } \theta$$
$$\theta^T x^{(i)} = p^{(i)} \|\theta\|$$
$$= \theta_1 x_1^{(i)} + \theta_2 x_2^{(i)}$$

Optimization Objective: Modified as follows:

$$\min_{\theta} \frac{1}{2} \sum_{j=1}^{n} \theta_j^2 = \frac{1}{2} \|\theta\|^2 \text{ s.t. } \begin{cases} p^{(i)} \|\theta\| \geq 1 & \text{if } y^{(i)} = 1 \\ p^{(i)} \|\theta\| \leq -1 & \text{if } y^{(i)} = 0 \end{cases}$$

where $p^{(i)}$ is the projection of $x^{(i)}$ on vector θ. $p^{(i)}$ will be bigger with large margin helping minimize the $\frac{1}{2} \|\theta\|^2$.

12.3 Kernels

12.3.1 Non-linear Decision Boundary

Predict $y = 1$ if $\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots \geq 0$.

New Notation: $h\theta = \theta_0 + \theta_1 f_1 + \theta_2 f_2 + \ldots$ where possibly, $f_1 = x_1$, $f_2 = x_2$, $f_3 = x_1 x_2$, $f_4 = x_1^2$, ….

Is there a better/different choice of features f_1, f_2, f_3, …?

Given x, choose 3 random landmarks ls and define new features as:

$$f_1 = similarity(x, l^{(1)}) \qquad = \exp\left(\frac{-\left\|x - l^{(1)}\right\|^2}{2\sigma^2}\right)$$

$$f_2 = similarity(x, l^{(2)}) \qquad = \exp\left(\frac{-\left\|x - l^{(2)}\right\|^2}{2\sigma^2}\right)$$

$$f_3 = similarity(x, l^{(3)}) \qquad = \exp\left(\frac{-\left\|x - l^{(3)}\right\|^2}{2\sigma^2}\right)$$

Similarity function is *Kernel* functions. This particular one is *Gaussian* Kernel.

Kernels and Similarity

If $x \approx l^{(1)}$: $f_1 \approx 1$
If x is far from $l^{(1)}$: $f_1 \approx 0$
Measures closeness to landmarks.
Setting σ^2 low, the value of f_1 falls to 0 very rapidly, and vice-versa.

Predict 1 when $\theta_0 + \theta_1 f_1 + \theta_2 f_2 + \theta_3 f_3 \geq 0$ Say $\theta_0 = -0.5$, $\theta_1 = 1$, $\theta_2 = 1$, $\theta_3 = 0$, then points close to $l^{(1)}$ or $l^{(2)}$ will give 1 creating complex non-linear boundary.

12.3.2   Choosing the Landmarks

Where to get landmarks? How many?

One choice would be one landmark per location of each feature. Given m examples, we have m landmarks.

Given an example x, compute f_1, f_2, … using kernels and get the feature vector f (include $f_0 = 1$). For training examples, at least one feature would be 1.

Hypothesis: Given x, compute $f \in \mathbb{R}^{m+1}$
Predict 'y=1' if $\theta^T f \geq 0$
Predict 'y=0' otherwise

Training:

$$\min_\theta C\left[\sum_{i=1}^m \left(y^{(i)} cost_1(\theta^T f^{(i)}) + (1 - y^{(i)}) cost_0(\theta^T f^{(i)})\right)\right] + \frac{1}{2}\sum_{j=1}^n \theta_j^2$$

$n = m$

Slight modification in last term! Instead of using $\theta^T \theta$, use $\theta^T M \theta$ where M is a matrix depending on the kernel done for computational expense.

Can use Kernels in other algorithms but tricks (such as M) don't go well with other algorithms so it can be slow/expensive.

12.3.3   SVM Parameters:

$C\left(= \dfrac{1}{\lambda}\right)$

$$\text{Large } C : \text{Lower Bias, High Variance}$$
$$\text{Small } C : \text{Higher Bias, Low Variance}$$

σ^2

$$\text{Large } \sigma^2 : \text{Features } f_i \text{ vary more smoothly}$$
$$: \text{Higher bias, Low Variance}$$
$$\text{Small } \sigma^2 : \text{Features } f_i \text{ vary less smoothly}$$
$$: \text{Lower Bias, High Variance}$$

12.4   Using SVM

Use SVM software package (e.g. liblinear, libsvm, ...) to solve for parameters θ.

Need to specify:

- Choice of parameter C

- Choice of kernel (similarity function):

    - Linear Kernel is just another name of no kernel.
    - Gaussian Kernel. Need to choose σ^2.

        ```
        function f = kernel(x1, x2)
        ```
        $$f = \exp\left(-\frac{\|x1 - x2\|^2}{2\sigma^2}\right)$$
        ```
        return
        ```

    Note: Do perform feature scaling before using the Gaussian Kernel.

– Other choices of kernel
  Note: Not all similarity functions make valid kernels. (Need to satisfy a technical condition called *Mercer's Theorem* to make sure SVM packages' optimizations run correctly, and do not diverge)
  Many off-the-shelf kernels available:

  * Polynomial Kernel: $k(x, l) = (x^T l + \text{constant})^{degree}$ (almost always performs worse than Gaussian Kernel, and not used that much and used only where x and l are all strictly positive)
  * More esoteric: String kernel, chi-square kernel, histogram intersection kernel, ...

12.4.1 Multi-Class Classification

Many SVM packages already have built-in multi-class classification functionality. Otherwise, use one vs all method.

12.4.2 Logistic Regression vs SVM

If $n \geq m$ (say $n = 10,000$ and $m = 10$ to 1000), use logistic regression or SVM without kernel (linear kernel).

If n is small (say upto 1000), m is intermediate (say upto $10,000$), use SVM with Gaussian Kernel.

If n is small (say upto 1000), m is large (say $50,000+$), SVM with Gaussian Kernels may start struggling. Create/Add more features, then use logistic regression or SVM without a kernel.

Neural Network likely to work well for most of these settings, but may be slower to train.

SVM is a convex optimization problem, so no need to worry about local optima.

Check Lecture12.pdf for more details.

Part II

Unsupervised Learning

Chapter 13

Clustering

13.1   Unsupervised Learning

Learning from unlabelled data as opposed to supervised learning. Give data to the algorithm to find some *structure* in our data.

13.2   Clustering

Group data in different *clusters.*

13.2.1   Application

- Market Segmentation
- Social Network Analysis
- Organize Computer Clusters
- Astronomical Data Analysis

13.3 K-means

13.3.1   Algorithm

Informally: Suppose we want to group the data into two clusters.

1. Randomly initialize two cluster centroids.

2. Assign each data point to a cluster centroid based on some *distance* measure.

3. Update cluster centroid to the average of each data point in the same cluster.

4. Go back to step 2 until convergence.

Formally :
Input:

- K (number of clusters)

- Training set $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$

Remark. $x^{(i)} \in \mathbb{R}^n$ (Drop $x_0 = 1$ as convention)

```
Randomly initialize K cluster centroids μ₁,μ₂,...,μₖ ∈ ℝⁿ
Repeat {
        for i = 1 to m
                c⁽ⁱ⁾ := index (from 1 to K) of cluster centroid
                        closest to x⁽ⁱ⁾
        for k = 1 to K
                μₖ := average (mean) of points assigned to cluster k
}
```

Remark. If we get a cluster centroid with zero data points in it, the common thing to do is to entirely eliminate that centroid and get $K-1$ clusters. But if you really need K clusters, you can randomly reinitialize the centroid.

13.3.2 K-means for non-separated clusters

Example: T-shirt sizing (S/M/L) based on weight and height.

13.3.3 Optimization Objective

Notations

$c^{(i)}$ = index of cluster $(1, 2, \dots, K)$ to which example $x^{(i)}$ is currently assigned

μ_k = cluster centroid k ($\mu_k \in \mathbb{R}^n$)

$\mu_{c^{(i)}}$ = cluster centroid of cluster to which example $x^{(i)}$ has been assigned

Cost Function

Also called *distortion* cost function or *distortion* of the K-means algorithm.

$$J(c^{(1)}, \ldots, c^{(m)}, \mu_1, \ldots, \mu_K) = \frac{1}{m} \sum_{i=1}^{m} \left\| x^{(i)} - \mu_{c^{(i)}} \right\|^2$$

Optimization Objective

$$\min_{\substack{c^{(1)}, \ldots, c^{(m)}, \\ \mu_1, \ldots, \mu_K}} J(c^{(1)}, \ldots, c^{(m)}, \mu_1, \ldots, \mu_K)$$

13.3.4 Random Initialization

Many ways, but this one usually works better in most cases. Should have $K < m$. Randomly pick K training examples as cluster centroids.

Local Optima

On bad days, K-means can get stuck in local optima. To avoid, we can initialize multiple times. Concretely, repeat about 100 times (usually 50 to 1000 times) and choose the way with least distortion.

For small K (2 to 10), multiple initialization is likely to give better solution while for larger Ks, it's highly likely that your first solution was pretty good already and not make a huge difference.

13.3.5 Choosing K

Still the most common thing is to choose the number of clusters by hand. Below are other thoughs.

Elbow Method

Plot distortion against the number of clusters. Find the *elbow* in the plot. Maybe that is a good number of clusters. Pretty good way if the graph goes down rapidly and then slowly later. But, not used that often coz often the elbow is very often very ambiguous.

Sometimes, we run K-means to get clusters for some later purpose. Evaluate based on metric for how well it performs for that later purpose.

Check Lecture13.pdf for more details.

Chapter 14

Dimensionality Reduction

14.1 Motivation

14.1.1 Data Compression

For example reduce data from 2D to 1D ($x^{(i)} \in \mathbb{R}^2 \to z^{(i)} \in \mathbb{R}$). This also helps to reduce features along with reducing memory requirement. More interestingly, it also helps speed up our learning algorithms.

14.1.2 Data Visualization

Helps us to visualize data in a better way since plotting a visualizing data in smaller dimensions is easier.

14.2 Principle Component Analysis

14.2.1 Problem Formulation

Remark. Before applying *PCA* perform mean normalization and feature scaling.

The goal of *PCA,* if we want to reduce data from 2D to 1D, is, we're going to find a vector in 2D to project the data minimizing the projection error.

More generally, to reduce from n-dimensions to k-dimensions, find k vectors $u^{(1)}, u^{(2)}, \ldots, u^{(k)}$ onto which to project the data, so as to minimize the projection error.

Remark. PCA is not *Linear Regression*

14.3 Algorithm

14.3.1 Data Preprocessing

Important! Always perform mean normalization and feature scaling.

14.3.2 Procedure

- Find component vectors $u^{(i)}$s

- Find reduced data values $z^{(i)}$s

Reduce data from n-dimensions to k-dimensions.

Compute *covariance matrix*:

$$\Sigma = \frac{1}{m} \sum_{i=1}^{m} (x^{(i)})(x^{(i)})^T$$

$$\text{if } X = \begin{bmatrix} -- & (x^{(1)})^T & -- \\ & \vdots & \\ -- & (x^{(m)})^T & -- \end{bmatrix}$$

$$\Sigma = \frac{1}{m} X^T X$$

Compute *eigenvectors* of matrix Σ: SVD stands for *Singular Value Decomposition*

```
[U,S,V] = svd(Σ);
```

There's also the `eig` function but `svd` is more numerically stable, though when you apply it on Σ, it'll give you the same values and that's because Σ satisfies *Symmetric Positive Semidefinite.*

Remark. $\Sigma \in \mathbb{R}^{n \times n}$

All we need from `svd` function is the $U \in \mathbb{R}^{n \times n}$ matrix.
The columns of matrix U will be exactly vectors $u^{(i)}$s we want.

$$U = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \dots & u^{(n)} \\ | & | & & | \end{bmatrix} \in \mathbb{R}^{n \times n}$$

Just take the first k vectors.

$$U_{reduce} = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \dots & u^{(k)} \\ | & | & & | \end{bmatrix} \in \mathbb{R}^{n \times k}$$

$$z = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \dots & u^{(k)} \\ | & | & & | \end{bmatrix}^T x \in \mathbb{R}^k$$

$$= \begin{bmatrix} -- & (u^{(1)})^T & -- \\ & \vdots & \\ -- & (u^{(k)})^T & -- \end{bmatrix} x \in \mathbb{R}^k$$

Full procedure in Octave will look like:

```
Σ = (1/m)*X'*X;
[U, S, V] = svd(Σ);
Ureduce = U(:, 1:k);
z = Ureduce'*x;
```

14.3.3 Reconstruction

$$z = U_{reduce}^T x$$
$$x_{approx} = U_{reduce} z$$

14.4 Choosing the number of principle components

$$\text{Average squared projection error} : \frac{1}{m} \sum_{i=1}^{m} \left\| x^{(i)} - x_{approx}^{(i)} \right\|^2$$

$$\text{Total variation in the data} : \frac{1}{m} \sum_{i=1}^{m} \left\| x^{(i)} \right\|^2$$

Typically, choose k to be the smallest value so that

$$\frac{\frac{1}{m} \sum_{i=1}^{m} \left\| x^{(i)} - x_{approx}^{(i)} \right\|^2}{\frac{1}{m} \sum_{i=1}^{m} \left\| x^{(i)} \right\|^2} \leq 0.01 \ (1\%)$$

99% of variance is retained

Other popular numbers are 95% or 90% variance retained.

14.4.1 Algorithm

One way is to try PCA with $k = 1, 2, \ldots$ and calculate variance retained till you get 99% or more. This procedure is horribly inefficient.

From [U, S, V] = svd(Σ), $S \in \mathbb{R}^{n \times n}$ is a diagonal matrix. For, a given k:

$$\frac{\frac{1}{m} \sum_{i=1}^{m} \left\| x^{(i)} - x_{approx}^{(i)} \right\|^2}{\frac{1}{m} \sum_{i=1}^{m} \left\| x^{(i)} \right\|^2} = 1 - \frac{\sum_{i=1}^{k} S_{ii}}{\sum_{i=1}^{n} S_{ii}}$$

So, we can test:

$$\frac{\sum_{i=1}^{k} S_{ii}}{\sum_{i=1}^{n} S_{ii}} \geq 0.99$$

This is much more efficient!

14.5 Applying PCA

14.5.1 Supervised Learning Speedup

$(x^{(i)}, y^{(i)})$ where say $x^{(i)} \in \mathbb{R}^{10000}$

1. Extract Inputs: Unlabelled dataset ($x^{(i)} \in \mathbb{R}^{10000}$)

2. Apply PCA (say 10x saving) ($x^{(i)} \to z^{(i)} \in \mathbb{R}^{1000}$)

3. New Training Set: $(z^{(i)}, y^{(i)})$

4. For test example x, map using same PCA to get z

Remark. Mapping $x^{(i)} \to z^{(i)}$ should be defined by running PCA only on the *Training Set*. This mapping can be applied as well to the examples in cross-validation and test sets.

14.5.2 Visualization

Choose $k = 2$ or $k = 3$ coz we know how to plot only in 2 or 3 dimensions.

14.5.3 Bad use of PCA: To prevent overfitting

Use $z^{(i)}$ instead of $x^{(i)}$ to reduce the number of features to $k < n$. Thus, fewer features, less likely to overfit.

This might work OK, but isn't a good way to address overfitting. Use regularization instead.

Bad because it throws away information.

14.5.4 PCA is sometimes used where it shouldn't be

Design of ML system:

1. Get training set

2. Run PCA to reduce dimensions (features)

3. Train

4. Test

Good question to ask: How about doing the whole thing without PCA?

Before implementing PCA, first try running whatever you want to do with the original/raw data. Only if that doesn't do what you want, then implement PCA.

Check Lecture14.pdf for more details.

Chapter 15

Anomaly Detection

15.1 Example

15.1.1 Aircraft Engine!

Features:

$$x_1 = \text{heat generated}$$
$$x_2 = \text{vibration intensity}$$
$$\ldots$$

Dataset: $\{x^{(1)}, x^{(2)}, \ldots, x^{(m)}\}$

Question: Is x_{test} anomalous?

Answer: Model $p(x)$ from dataset, and if $p(x_{test}) < \epsilon$ then flag anomaly otherwise OK!

15.1.2 Fraud Detection

Identify unusual users by checking $p(x) < \epsilon$

15.1.3 Manufacturing

Just like Aircraft Engine

15.1.4 Monitoring computers in a data center

Check which system is probably required a review by a system administrator.

15.2 Gaussian (Normal) Distribution

15.2.1 Definition

Say $x \in \mathbb{R}$. If x is distributed *Gaussian* with mean μ and variance σ^2 (σ is the standard deviation).

$$x \sim \mathcal{N}(\mu, \sigma^2)$$

$$p(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

Bell shaped curve centered at μ and width varying by σ

15.2.2 Parameter Estimation

Given a dataset, estimate μ and σ or σ^2

Maximum Likelyhood Estimate (MLE):

$$\mu = \frac{1}{m} \sum_{i=1}^{m} x^{(i)}$$

$$\sigma^2 = \frac{1}{m} \sum_{i=1}^{m} (x^{(i)} - \mu)^2$$

15.3 Algorithm

15.3.1 Density Estimation

Training Set: $\{x^{(1)}, \ldots, x^{(m)}\}$
Each example is $x \in \mathbb{R}$
Model $p(x)$ as such:

$$p(x) = p(x_1)p(x_2)\ldots p(x_n)$$

Assume each feature is distributed Gaussian independently

$$x_1 \sim \mathcal{N}(\mu_1, \sigma_1^2)$$
$$x_2 \sim \mathcal{N}(\mu_2, \sigma_2^2)$$
$$\vdots$$
$$x_n \sim \mathcal{N}(\mu_n, \sigma_n^2)$$
$$p(x) = p(x_1; \mu_1, \sigma_1^2) p(x_2; \mu_2, \sigma_2^2) \dots p(x_n; \mu_n, \sigma_n^2)$$
$$= \prod_{j=1}^{n} p(x_j; \mu_j, \sigma_j^2)$$

In practice, it works fine even if the features are not really *independent*.

15.3.2   Anomaly Detection Algorithm

1. Choose features x_i that might be indicative of anomaly.

2. Given a training set $\{x^{(1)}, \dots, x^{(m)}\}$, fit paramters $\mu_1, \dots, \mu_n, \sigma_1^2, \dots, \sigma_2^2$

$$\mu = \frac{1}{m} \sum_{i=1}^{m} x^{(i)}$$
$$\sigma^2 = \frac{1}{m} \sum_{i=1}^{m} (x^{(i)} - \mu)^2$$

3. Given a new example x, compute $p(x)$

$$p(x) = \prod_{j=1}^{n} p(x_j; \mu_j, \sigma_j^2)$$
$$= \prod_{j=1}^{n} \frac{1}{\sqrt{2\pi}\sigma_j} \exp\left(-\frac{(x_j - \mu_j)^2}{2\sigma_j^2}\right)$$

4. Anomaly if $p(x) < \epsilon$

15.4 Developing and Evaluating an Anomaly Detection System

15.4.1 The importance of real-number evaluation

When developing a learning algorithm (choosing features, etc.), making decisions is much easier if we have a way of evaluating our learning algorithm.

Assume we have some labelled data, of anomalous and non-anomalous examples. ($y = 0$ if normal, $y = 1$ if anomalous).

Training set: $x^{(1)}, x^{(2)}, \ldots, x^{(m)}$ (assume normal)

Cross validation set: $(x_{cv}^{(1)}, y_{cv}^{(1)}), \ldots, (x_{cv}^{(m_{cv})}, y_{cv}^{(m_{cv})})$

Test set: $(x_{test}^{(1)}, y_{test}^{(1)}), \ldots, (x_{test}^{(m_{test})}, y_{test}^{(m_{test})})$

Assume we include examples in cross validation set and test set which are known to be anomalous.

Aircraft Engines Example

Recommended split:

- 10000 good (normal) engines

- 20 - 50 flawed (anomalous) engines

- Training Set: 6000 good engines

- CV: 2000 good engines, 10 anomalous

- Test: 2000 good engines, 10 anomalous

15.4.2 Algorithm Evaluation

1. Fit model $p(x)$ on training set $\{x^{(1)}, \ldots, x^{(m)}\}$

2. On a cross-validation/test example x, predict

$$y = \begin{cases} 1 & \text{if } p(x) < \epsilon \text{ (anomaly)} \\ 0 & \text{otherwise} \end{cases}$$

3. Possible evaluation metrics

- True positive, false positive, false negative, true negative
- Precision/Recall
- F_1-score

4. Can also use cross validation set to choose parameter ϵ

15.5 Anomaly Detection vs Supervised Learning

If we have labelled data, why not use supervised learning?

Anomaly Detection	Supervised Learning
Very small number of positive examples ($y = 1$). Large number of negative examples ($y = 0$).	Large number of positive and negative examples.
Many different *types* of anomalies. Hard for any algorithm to learn from positive examples what the anomalies look like.	Enough positive examples for algorithm to get a sense of what positive examples are like.
Future anomalies may look nothing like any of the anomaly examples we've seen so far.	Future positive examples likely to be similar to ones in training set.
Fraud Detection	Email spam classification
Manufacturing (e.g. aircraft engine)	Weather prediction (sunny/rainy/etc.)
Monitoring machines in a data center	Cancer classification
$\vdots$	$\vdots$

15.6 Choosing what features to use

15.6.1 Transformations

Works fine even if data isn't gaussian, but usually is a good sanity check for a feature.

If not gaussian, play with different kinds of transformations and get something similar to gaussian, like log transformation.

Other transformations:

$$x_1 \leftarrow \log x_1$$
$$x_1 \leftarrow \log x_2 + c$$
$$x_3 \leftarrow \sqrt{x_3} = x_3^{\frac{1}{2}}$$
$$x_4 \leftarrow x_4^{\frac{1}{3}}$$

15.6.2   Error analysis for anomaly detection

Want $p(x)$ large for normal examples x.

$p(x)$ small for anomalous examples x.

Most common problem: $p(x)$ is comparable (say, both large) for normal and anomalous examples.

Look at such anomalous examples and try to get a new feature to distinguish.

15.6.3   Other techniques:

Choose features that might take on unusually large or small values in the event of an anomaly.

Example: Monitoring computers in a data center running servers.

$$x_1 = \text{memory use}$$
$$x_2 = \text{disk access}$$
$$x_3 = \text{CPU load}$$
$$x_4 = \text{network traffic}$$

Possible anomaly = infinite loop, so $x_5 = \dfrac{\text{CPU Load}}{\text{network traffic}}$

15.7   Multivariate Gaussian (Normal) Distribution

Example

Monitoring machines in a data center. Consider $x_1 = $ CPU load and $x_2 = $ Memory Use. An outlier in 2D plot would not be triggered anomalous when treating x_1 and x_2.

$x \in \mathbb{R}^n.$ Don't model $p(x_1), p(x_2), \ldots$, etc. separately. Model $p(x)$ all in one go.

Parameters: $\mu \in \mathbb{R}^n, \Sigma \in \mathbb{R}^{n \times n}$ (covariance matrix)

Equation:

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right)$$

where $|\Sigma| = $ determinant of $\Sigma = $ `det`(Σ) (in octave)

15.8 Anomaly Detection using Multivariate Gaussian Distribution

15.8.1 Model

Parameters: μ, Σ

Parameter Fitting: Given training set $\{x^{(1)}, \ldots, x^{(m)}\}$

$$\mu = \frac{1}{m} \sum_{i=1}^{m} x^{(i)}$$

$$\Sigma = \frac{1}{m} \sum_{i=1}^{m} \left(x^{(i)} - \mu\right)\left(x^{(i)} - \mu\right)^T$$

15.8.2 Anomaly Detection

1. Fit model $p(x)$ by setting μ, Σ

2. Given a new example x, compute

$$p(x) = \frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right)$$

3. Flag an anomaly if $p(x) < \epsilon$

15.8.3 Relationship to original model

Original model corresponds to multivariate gaussian where the contours of the gaussian are axis aligned. The contraint (mathematically) is

$$\Sigma = \begin{bmatrix} \sigma_1^2 & 0 & 0 & \ldots & 0 \\ 0 & \sigma_2^2 & 0 & \ldots & 0 \\ 0 & 0 & \ddots & & 0 \\ 0 & \ldots & 0 & \sigma_{n-1}^2 & 0 \\ 0 & \ldots & 0 & 0 & \sigma_n^2 \end{bmatrix}$$

Original Model	Multivariate Gaussian
$\displaystyle\prod_{j=1}^{n} p(x_j; \mu_j, \sigma_j^2)$	$\dfrac{1}{(2\pi)^{\frac{n}{2}} \lvert\Sigma\rvert^{\frac{1}{2}}} e^{\left(-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)\right)}$
Manually create features to capture anomalies where x_1, x_2 take unusual combination of values.	Automatically captures correlations between features.
Computationally cheaper (scales better to large n, say 100,000).	Computationally more expensive.
OK even if m (training set size) is small.	Must have $m > n$, or else Σ is non-invertible. (Typical rule of thumb: $m \geq 10n$)

Remark. If Σ is singular (non-invertible), either $m > n$ fails, or you have redundant (linearly dependent) features.

Check Lecture15.pdf for more details.

Chapter 16

Recommender Systems

16.1 Problem Formulation

Example: Predicting movie ratings! Collect movie ratings from users for the movies they have watched, and predict ratings for the movies they haven't watched.

Notations:

$$
\begin{aligned}
n_u &= \text{no. users} \\
n_m &= \text{no. movies} \\
r(i,j) &= 1 \text{ if user } j \text{ has rated movie } i \\
y^{(i,j)} &= \text{rating given by user } j \text{ to movie } i \\
&\quad (\text{defined only if } r(i,j) = 1)
\end{aligned}
$$

16.2 Content-Based Recommendations

Features could be the degree of genre of the movie, say romance or action.

One way could be, for each user j, learn a parameter $\theta^{(j)} \in R^{n+1}$. Predict user j as rating movie i with $(\theta^{(j)})^T x^{(i)}$ stars.

Notations:

$$n = \text{no. features}$$
$$\theta^{(j)} = \text{parameter vector for user } j$$
$$x^{(i)} = \text{feature vector for movie } i$$
$$(\theta^{(j)})^T x^{(i)} = \text{predicted rating for user } j, \text{ movie } i$$
$$m^{(j)} = \text{no. movies rated by user } j$$

To learn $\theta^{(j)}$:

$$\min_{\theta^{(j)}} \left(\frac{1}{2m^{(j)}} \sum_{i:r(i,j)=1} \left(\left(\theta^{(j)}\right)^T \left(x^{(i)}\right) - y^{(i,j)} \right)^2 + \frac{\lambda}{2m^{(j)}} \sum_{k=1}^{n} \left(\theta_k^{(j)}\right)^2 \right)$$

For recommender systems, we make some simplifications:

- Get rid of the constant $m^{(j)}$

16.2.1 Optimization Objective

To learn $\theta^{(j)}$ (parameter for user j):

$$\min_{\theta^{(j)}} J(\theta^{(j)})$$

$$\min_{\theta^{(j)}} \left(\frac{1}{2} \sum_{i:r(i,j)=1} \left(\left(\theta^{(j)}\right)^T \left(x^{(i)}\right) - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{k=1}^{n} \left(\theta_k^{(j)}\right)^2 \right)$$

To learn $\theta^{(1)}, \theta^{(2)}, \ldots, \theta^{(n_u)}$:

$$\min_{\theta^{(1)}, \ldots, \theta^{(n_u)}} J(\theta^{(1)}, \ldots, \theta^{(n_u)})$$

$$\min_{\theta^{(1)}, \ldots, \theta^{(n_u)}} \left(\frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} \left(\left(\theta^{(j)}\right)^T \left(x^{(i)}\right) - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^{n} \left(\theta_k^{(j)}\right)^2 \right)$$

16.2.2 Optimization Algorithm

Gradient descent update:

$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \sum_{i:r(i,j)=1} \left(\left(\theta^{(j)}\right)^T x^{(i)} - y^{(i,j)} \right) x_k^{(i)} \qquad \text{(for } k = 0)$$

$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \left(\sum_{i:r(i,j)=1} \left(\left(\theta^{(j)}\right)^T x^{(i)} - y^{(i,j)} \right) x_k^{(i)} + \lambda \theta_k^{(j)} \right) \qquad \text{(for } k = 0)$$

16.3 Collaborative Filtering

Feature Learning!

Each user j just tells us how much they like different types of movies $\theta^{(j)}$. If we get these parameters from the users, then it is possible to infer what are the features for each movie from the ratings.

Optimization Algorithm

Given $\theta^{(1)}, \theta^{(2)}, \ldots, \theta^{(n_u)}$, to learn $x^{(i)}$:

$$\min_{x^{(i)}} \left(\frac{1}{2} \sum_{j:r(i,j)=1} \left(\left(\theta^{(j)} \right)^T \left(x^{(i)} \right) - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{k=1}^{n} \left(x_k^{(i)} \right)^2 \right)$$

Given $\theta^{(1)}, \theta^{(2)}, \ldots, \theta^{(n_u)}$, to learn $x^{(1)}, \ldots, x^{(n_m)}$:

$$\min_{x^{(i)}} \left(\frac{1}{2} \sum_{i=1}^{n_m} \sum_{j:r(i,j)=1} \left(\left(\theta^{(j)} \right)^T \left(x^{(i)} \right) - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^{n} \left(x_k^{(i)} \right)^2 \right)$$

Collaboration

Guess $\theta \to x \to \theta \to x \to \ldots$

16.3.1 Algorithm

Minimizing $x^{(1)}, \ldots, x^{(n_m)}$ and $\theta^{(1)}, \ldots, \theta^{(n_u)}$ simultaneously:

$$J(x^{(1)}, \ldots, x^{(n_m)}, \theta^{(1)}, \ldots, \theta^{(n_u)}) = \frac{1}{2} \sum_{(i,j):r(i,j)=1} \left(\left(\theta^{(j)} \right)^T x^{(i)} - y^{(i,j)} \right)^2$$

$$+ \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^{n} \left(x_k^{(i)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^{n} \left(\theta^{(j)} \right)^2$$

$$\min_{\substack{x^{(1)}, \ldots, x^{(n_m)} \\ \theta^{(1)}, \ldots, \theta^{(n_u)}}} J(x^{(1)}, \ldots, x^{(n_m)}, \theta^{(1)}, \ldots, \theta^{(n_u)})$$

When learning features this way, we're going to do away with $x_0 = 1$ and hence, now our $x \in \mathbb{R}^n$ and $\theta \in \mathbb{R}^n$, and the reason is that now that we're learning the features, if there is a need for such a feature, the algorithm can learn it by itself.

Collaborative Filtering Algorithm

1. Initialize $x^{(1)},\ldots,x^{(n_m)}$ and $\theta^{(1)},\ldots,\theta^{(n_u)}$ to small random values (just like neural network).

2. Minimize $J(x^{(1)},\ldots,x^{(n_m)},\theta^{(1)},\ldots,\theta^{(n_u)})$ using Gradient Descent (or an advanced optimization algorithm).
E.g. for every $j = 1,\ldots,n_u, i = 1,\ldots,n_m$:

$$x_k^{(i)} := x_k^{(j)} - \alpha \left(\sum_{i:r(i,j)=1} \left(\left(\theta^{(j)}\right)^T x^{(i)} - y^{(i,j)} \right) \theta_k^{(j)} + \lambda x_k^{(i)} \right)$$

$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \left(\sum_{i:r(i,j)=1} \left(\left(\theta^{(j)}\right)^T x^{(i)} - y^{(i,j)} \right) x_k^{(i)} + \lambda \theta_k^{(j)} \right)$$

3. For a user with parameters θ and a movie with (learned) features x, predict a star rating of $\theta^T x$.

16.4 Vectorization: Low Rank Matrix Vectorization

16.4.1 Vectorization

$$Y = \left[y^{(i,j)} \right]$$
$$= \left[\left(\theta^{(j)}\right)^T x^{(i)} \right]$$

$$X = \begin{bmatrix} \text{---} & \left(x^{(1)}\right)^T & \text{---} \\ \text{---} & \left(x^{(2)}\right)^T & \text{---} \\ & \vdots & \\ \text{---} & \left(x^{(n_m)}\right)^T & \text{---} \end{bmatrix}$$

$$\Theta = \begin{bmatrix} \text{---} & \left(\theta^{(1)}\right)^T & \text{---} \\ \text{---} & \left(\theta^{(2)}\right)^T & \text{---} \\ & \vdots & \\ \text{---} & \left(\theta^{(n_u)}\right)^T & \text{---} \end{bmatrix}$$

$$Y = X\Theta^T$$

Name *Low Rank Matrix Vectorization* comes from the fact that $X\Theta^T$ is a *Low Rank Matrix*.

16.4.2 Finding related movies

Q: For each product i, we learn a feature vector $x^{(i)} \in \mathbb{R}^n$. How to find movies j related to movie i?

A: Small $\left\| x^{(i)} - x^{(j)} \right\| \to$ movies j and i are similar.
5 most similar movies to i:
Find the 5 movies j with the smallest $\left\| x^{(i)} - x^{(j)} \right\|$

16.5 Implementational Detail: Mean Normalization

Let's consider an example of a user who hasn't rated any movie. So in our optimization objective

$$
\min_{\substack{x^{(1)},\ldots,x^{(n_m)} \\ \theta^{(1)},\ldots,\theta^{(n_u)}}} \frac{1}{2} \sum_{(i,j):r(i,j)=1} \left(\left(\theta^{(j)}\right)^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^{n} \left(x_k^{(i)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^{n} \left(\theta^{(j)} \right)^2
$$

the first term plays no role, and hence we have to minimize $\frac{\lambda}{2}(\theta_1^2 + \theta_2^2)$ which will give $\theta = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$ and so we will predict $\theta^T x = 0$ for all x, which will give us no good way to recommend movies to the new user. So, calculate mean rating for each movie and subtract the mean from all ratings (so that each movie has a mean rating of zero).

$$
Y = \begin{bmatrix} y^{(1,j)} \\ y^{(2,j)} \\ \vdots \\ y^{(n_m,j)} \end{bmatrix}
$$

$$
\mu = \begin{bmatrix} \text{mean}(y^{(1)}) \\ \text{mean}(y^{(2)}) \\ \vdots \\ \text{mean}(y^{(n_m)}) \end{bmatrix}
$$

$$
Y_{new} := Y - \mu
$$

Now, learn $\theta^{(j)}, x^{(i)}$ on this Y_{new}.
For user j, on the movie i predict: $(\theta^{(j)})^T(x^{(i)}) + \mu_i$

Check Lecture16.pdf for more details.